

DEEP LEARNING MODEL ON MNIST DATASET

The architecture of this model comprises of three hidden layer with 500, 250 and 100 being their respective sizes

The aim is to model a complete handwritten digit recognizer, train it on the train set comprising 60,000 images and test performance on the test set comprising 10,000 images

Baseline model includes:

- Training data size of MNIST data = 60,000
 - with its respective batch size being 64
- Testing data size of MNIST data = 10,000
- Learning rate = 0.01
- number of epochs = 15

Importing Libraries

```
In [1]: 1 import numpy as np
        2 import matplotlib.pyplot as plt
        3 import torch
        4 import torchvision.datasets as data
        5 from torchvision.transforms import ToTensor
        6 from torch.utils.data import DataLoader
        7 from sklearn.metrics import confusion_matrix, classification_report
        8
        9 %matplotlib inline
       10 plt.rcParams['figure.figsize'] = (7.0, 4.0) # set default size of plots
       11
```

Setting up to GPU

```
In [2]: 1 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
        2 device
```

```
Out[2]: device(type='cuda', index=0)
```

Importing the MNIST dataset


```
In [3]: 1 train_set = data.MNIST(root = 'MNIST/raw/train-images-idx3-ubyte', train = True, transform= ToTensor(), downl
        2 test_set = data.MNIST(root= 'MNIST/raw/train-images-idx3-ubyte', train= False, transform= ToTensor(), downlo
```

Loading the data using DataLoader

```
In [4]: 1 train_features = train_set.data #getting images
2 train_labels = train_set.targets #getting labels
3 # Visualising data
4 # Display image and label
5
6 print(f"Feature batch shape: {train_features.size()}")
7 print(f"Labels batch shape: {train_labels.size()}")
8
9 fig = plt.figure()
10 for i in range(6):
11     plt.subplot(2,3,i+1)
12     plt.tight_layout()
13     plt.imshow(train_features[i].squeeze(), cmap='gray', interpolation='none')
14     plt.title("Ground Truth: {}".format(train_labels[i]))
15     plt.xticks([])
16     plt.yticks([])
17 plt.show()
18
```

Feature batch shape: torch.Size([60000, 28, 28])

Labels batch shape: torch.Size([60000])



One Hot Encoding

```
In [5]: 1 """
2 convert the labels into one-hot format before training and testing the network
3 """
4 def one_hot_encode(Y):
5     output = np.eye(10)[np.array(Y).reshape(-1)]
6     return output.reshape(list(np.shape(Y))+[10])
7
```

Activation Functions

Actvation functions and their respective derivatives

(1) Sigmoid

(2) ReLu

(3) TanH

(2) Softmax

```
In [6]: 1 #ACTIVATION FUNCTIONS
2 def activation(x, a_type = 'SOFTMAX'):
3
4     if a_type == 'SIGMOID':
5         return 1/(1+np.exp(-x))
6
7     if a_type == 'RELU':
8         return np.maximum(0,x)
9
10    if a_type == 'TANH':
11        return ((np.exp(x)-np.exp(-x))/(np.exp(x)+np.exp(-x)))
12
13    if a_type == 'SOFTMAX':
14        return np.exp(x)/sum(np.exp(x))
15
16    #DERIVATIVES OF ACTIVATION FUNCTIONS FOR BACK PROPAGATION
17
18    def derivative(x, d_type = 'SIGMOID'):
19
20        if d_type == 'SIGMOID':
21            s = activation(x, 'SIGMOID')
22            return (s*(1-s))
23
24        elif d_type == 'RELU':
25            return (x>0)
26
27        elif d_type == 'TANH':
28            t = activation(x, 'TANH')
29            return (1-t**2)
```

BASELINE

(1) Data Flattening

```
In [7]: 1 # Flattening and normalising
2
3 def data_flattening(features, labels, one_hot = True):
4     features = features.numpy()
5     labels = labels.numpy()
6     X = (features.reshape(features.shape[0], -1)).T
7     if one_hot:
8         Y = one_hot_encode(labels)
9     else:
10        Y = labels
11    return X, Y
```

(2) Initialising Parameters

```
In [8]: 1 batchsize = 64
2 input_size = (train_features.reshape(train_features.shape[0], -1)).shape[1]
3 hidden_layer_1 = 500
4 hidden_layer_2 = 250
5 hidden_layer_3 = 100
6 output_layer = 10
7 n_epochs = 15
8 learning_rate = 0.01
```

```

In [9]: 1 """
2 Function to Initialise the parameters -weights and biases
3 W1, W2, W3, W4, b1, b2, b3, b4
4 for the network {i/p -> 500 -> 250 -> 100 ->o/p}
5 The parameters here are initialised by Xavier initialisation and is one of the common ways of initialisation
6 """
7 def initialise_parameter(dim):
8     np.random.seed(11)
9
10     parameters = {}
11     L = len(dim)
12     for i in range(1, L):
13         Ni = dim[i-1]
14         No = dim[i]
15         M = np.sqrt(6/(Ni+No))
16         parameters["W" + str(i)] = np.asarray(np.random.uniform(-M, M, size = (No, Ni)))
17         parameters["b" + str(i)] = np.zeros((dim[i], 1))
18
19         assert(parameters["W" + str(i)].shape == (dim[i], dim[i-1]))
20         assert(parameters["b" + str(i)].shape == (dim[i], 1))
21     return parameters
22

```

(3) Forward propagation function

- using Sigmoid
- using TanH
- using ReLu

```

In [10]: 1 #Function to implement forward propagation
          2 """
          3 Retrieving the parameters
          4 Linear function on the inputs followed by activation functions
          5
          6 ---
          7
          8 In this code we have only got the outputs for Sigmoid and TanH,
          9 for relu seperate code follows
         10
         11 """
         12 def forward_propagation(X, parameters, forward):
         13     # retrieve parameters
         14     W1 = parameters["W1"]
         15     b1 = parameters["b1"]
         16
         17     W2 = parameters["W2"]
         18     b2 = parameters["b2"]
         19
         20     W3 = parameters["W3"]
         21     b3 = parameters["b3"]
         22
         23     W4 = parameters["W4"]
         24     b4 = parameters["b4"]
         25
         26     # FORWARD PROPOGATION : LINEAR -> SIGMOID -> LINEAR -> SIGMOID -> LINEAR -> SIGMOID -> LINEAR -> SOFTMAX
         27     # FORWARD PROPOGATION : LINEAR -> TANH -> LINEAR -> TANH -> LINEAR -> TANH -> LINEAR -> SOFTMAX
         28     # FORWARD PROPOGATION : LINEAR -> RELU -> LINEAR -> RELU -> LINEAR -> RELU -> LINEAR -> SOFTMAX
         29
         30     Z1 = np.dot(W1, X) + b1
         31     A1 = activation(Z1, a_type = forward) # Sigmoid or ReLu or TanH
         32
         33     Z2 = np.dot(W2, A1) + b2
         34     A2 = activation(Z2, a_type = forward) # Sigmoid or ReLu or TanH
         35
         36     Z3 = np.dot(W3, A2) + b3
         37     A3 = activation(Z3, a_type = forward) # Sigmoid or ReLu or TanH
         38
         39     Z4 = np.dot(W4, A3) + b4
         40     A4 = activation(Z4, a_type = 'SOFTMAX') # Softmax
         41

```



```
42     cache = {  
43         "z1" : z1, "z2" : z2, "z3" : z3, "z4" : z4,  
44         "A1" : A1, "A2" : A2, "A3" : A3, "A4" : A4,  
45         "W1" : W1, "W2" : W2, "W3" : W3, "W4" : W4,  
46         "b1" : b1, "b2" : b2, "b3" : b3, "b4" : b4}  
47  
48     return A4, cache  
49
```

(4) Backward Propagation

```

In [11]: 1 #Function to implement Backward Propagation
          2 """
          3 Back Propagation is to get the gradients which will be used
          4 to update the parameters by gradient descent method
          5 """
          6 """
          7 def backward_propagation(X, Y, cache, activation):
          8
          9     m = batchsize
         10     A4 = cache["A4"]
         11     A3 = cache["A3"]
         12     A2 = cache["A2"]
         13     A1 = cache["A1"]
         14     Z4 = cache["Z4"]
         15     Z3 = cache["Z3"]
         16     Z2 = cache["Z2"]
         17     Z1 = cache["Z1"]
         18     W4 = cache["W4"]
         19     W3 = cache["W3"]
         20     W2 = cache["W2"]
         21     W1 = cache["W1"]
         22     b4 = cache["b4"]
         23     b3 = cache["b3"]
         24     b2 = cache["b2"]
         25     b1 = cache["b1"]
         26
         27     dZ4 = A4 - Y
         28     dW4 = 1./m * np.dot(dZ4, A3.T)
         29     db4 = 1./m * np.sum(dZ4, axis=1, keepdims = True)
         30
         31     dA3 = np.dot(W4.T, dZ4)
         32     dZ3 = np.multiply(dA3, derivative(A3, d_type = activation))
         33     dW3 = 1./m * np.dot(dZ3, A2.T)
         34     db3 = 1./m * np.sum(dZ3, axis=1, keepdims = True)
         35
         36     dA2 = np.dot(W3.T, dZ3)
         37     dZ2 = np.multiply(dA2, derivative(A2, d_type = activation))
         38     dW2 = 1./m * (np.dot(dZ2, A1.T))
         39     db2 = 1./m * np.sum(dZ2, axis=1, keepdims = True)
         40
         41     dA1 = np.dot(W2.T, dZ2)

```

```
42 dZ1 = np.multiply(dA1, derivative(A1, d_type = activation))
43 dW1 = 1./m * (np.dot(dZ1, X.T))
44 db1 = 1./m * np.sum(dZ1, axis=1, keepdims = True)
45
46 gradients = {"dW4": dW4, "db4": db4,
47             "dW3": dW3, "db3": db3,
48             "dW2": dW2, "db2": db2,
49             "dW1": dW1, "db1": db1}
50
51 return gradients
52
```

(5) Update Parameters

```
In [12]: 1 #Function to update the parameters
2         """
3         Updating the parameters for the next iteration
4         in order to train the model
5
6         ---
7
8         returns the updated gradients
9         """
10        def update_parameters(parameters, grads, learning_rate):
11            W1 = parameters["W1"]
12            b1 = parameters["b1"]
13            W2 = parameters["W2"]
14            b2 = parameters["b2"]
15            W3 = parameters["W3"]
16            b3 = parameters["b3"]
17            W4 = parameters["W4"]
18            b4 = parameters["b4"]
19            dW1 = grads["dW1"]
20            db1 = grads["db1"]
21            dW2 = grads["dW2"]
22            db2 = grads["db2"]
23            dW3 = grads["dW3"]
24            db3 = grads["db3"]
25            dW4 = grads["dW4"]
26            db4 = grads["db4"]
27            W1 = W1 - learning_rate * dW1
28            b1 = b1 - learning_rate * db1
29            W2 = W2 - learning_rate * dW2
30            b2 = b2 - learning_rate * db2
31            W3 = W3 - learning_rate * dW3
32            b3 = b3 - learning_rate * db3
33            W4 = W4 - learning_rate * dW4
34            b4 = b4 - learning_rate * db4
35
36            parameters={"W1":W1, "b1":b1,
37                        "W2":W2, "b2":b2,
38                        "W3":W3, "b3":b3,
39                        "W4":W4, "b4":b4}
40            return parameters
```

(6) Compute Cost

```
In [13]: 1 """
2 Function to compute cross entropy cost
3 ---
4 Cross Entropy Loss is used to compute the cost in case of classification models
5 """
6 def compute_cost(A, Y):
7     m = Y.shape[1]
8     logprobs = np.multiply(-np.log(A), Y) + np.multiply(-np.log(1 - A), (1 - Y))
9     cost = 1./m * np.nansum(logprobs)
10    return cost
```

(8) Predict Labels

```
In [14]: 1 # Predict
2 def predict(X, Y, parameters, batchsize = batchsize):
3
4     m = batchsize
5     y = Y.T
6     p = np.zeros((m), dtype = np.int)
7     y = y.numpy()
8     # Forward propagation
9     a4, caches = forward_propagation(X, parameters)
10    print(a4)
11    p = np.argmax(a4, axis = 0)
12    # print results
13    print(f"predicted_value is {p}")
14    print(f"label value is {y}")
15    print("Accuracy: " + str(np.mean((p == y))))
16    return p
```

(9) Accuracy

```
In [15]: 1 """
2 Computing the accuracy of the testing and training sets
3
4 ---
5 returns the actual and predicted value of y
6 prints the accuracy
7
8 """
9
10 def Accuracy(dataset, parameters, forward, size):
11
12     features = dataset.data
13     labels = dataset.targets
14     X, Y = data_flattening(features, labels, one_hot = False)
15     y = Y.T
16     p = np.zeros(size, dtype = int)
17     # Forward propagation
18     a4, caches = forward_propagation(X, parameters, forward)
19     p = np.argmax(a4, axis = 0)
20     a = np.mean((p == y))
21     print("accuracy is =" + str(a))
22     return y, p
```

(10) Gradient Descent Model

```

In [24]: 1 """
2 Model to implement Neural Network by Gradient Descent
3 ---
4 This function will run once for the entire dataset
5 ---
6 Input parameters are - train_loader, parameters, type of activation, Learning rate, iterations and lamda
7 ---
8 This function trains model for both regularised and unregularised
9
10 """
11 def gradient_descent(dataloader, parameters, forward, learning_rate = 0.01, number_of_iterations = int(60000)):
12
13     grads = {}
14     costs = []
15     cost_out = []
16     m = batchsize #number of examples
17     for i in range(0, number_of_iterations):
18         features, labels = next(iter(dataloader))
19         X, Y = data_flattening(features, labels, True)
20         a4, cache = forward_propagation(X, parameters, forward) # FORWARD PROPOGATION : LINEAR -> SIGMOID
21         # Cost function
22         if lambd == 0:
23             cost = compute_cost(a4, Y.T)
24             grads = backward_propagation(X, Y.T, cache, forward)
25         else:
26             cost = compute_cost_with_regularization(a4, Y.T, parameters, lambd)
27             grads = backward_propagation_with_regularization(X, Y.T, cache, lambd, forward)
28
29         parameters = update_parameters(parameters, grads, learning_rate) # Update parameters
30         costs.append(cost)
31         if i%200==0 or i==(number_of_iterations-1):
32             print("Cost after iteration {}: {}".format(i, cost))
33             #plot the cost
34             plt.plot(costs)
35             plt.ylabel('cost')
36             plt.xlabel('iterations')
37             plt.title("Learning rate =" + str(learning_rate))
38             plt.show()
39             costs = []
40         if i%500 == 0:
41             cost_out.append(cost)

```

```
42     acc = []
43     # Load images to a Torch Variable
44     images = train_set.data
45     labels_1 = train_set.targets
46     X, Y = data_flattening(images, labels_1, one_hot = False)
47     y = Y.T
48     predicted = np.zeros(60000, dtype = int)
49     # Forward propagation
50     a4, caches = forward_propagation(X, parameters, forward)
51     predicted = np.argmax(a4, axis = 0)
52     accuracy = np.mean((predicted == y))
53     # Print Loss
54     acc.append(accuracy)
55
56     return parameters, cost_out, acc
57
```

Neural Network


```

In [27]: 1 """
2 Neural Network Function
3 ---
4 Calls the model for the epoch number of times
5 """
6
7 def NeuralNet(dataset, forward = 'SIGMOID', learning_rate = 0.01, n_epochs = 1, batch_size = batchsize, lambda):
8     costs = []
9     accuracy = []
10    m = batch_size #number of examples
11    layers_dim = [input_size, hidden_layer_1, hidden_layer_2, hidden_layer_3, output_layer]
12    parameters = initialise_parameter(layers_dim)
13    for epoch in range(n_epochs):
14        print(f"Epoch :- {epoch+1}")
15        train_dataloader = DataLoader(train_set, batch_size=batchsize, shuffle=True)
16        parameters, cost, acc = gradient_descent(train_dataloader, parameters, forward, learning_rate = learning_rate)
17        print (f'Epoch [{epoch+1}/{n_epochs}], Cost: {cost[-1]}, Accuracy: {acc[-1]}')
18        accuracy.append(acc)
19        costs.append(cost)
20    #plot the cost
21    plt.plot(costs)
22    plt.ylabel('cost')
23    plt.xlabel('epochs')
24    plt.title("Learning rate =" + str(learning_rate))
25    plt.show()
26    return parameters, accuracy

```


Training

Sigmoid Activation Function

```
In [48]: 1 """
          2 Training the Neural Network
          3 Getting trained parameters and accuracy of train dataset
          4
          5 """
          6 parameters, train_acc = NeuralNet(train_set, forward = 'SIGMOID', learning_rate = 0.01, n_epochs = 15, batch_
          7
```

Epoch :- 1

Cost after iteration 0: 3.4647339306614118



```
In [49]: 1 print ("On the TEST set:")
          2 y_actual_test, y_pred_test = Accuracy(test_set, parameters, 'SIGMOID', 10000)
```

On the TEST set:
accuracy is =0.729


/home/mansi/.conda/envs/hbp/lib/python3.7/site-packages/ipykernel_launcher.py:5: RuntimeWarning: overflow encountered in exp
"""

```

In [52]: 1 #to get the heatmap for the confusion matrix
2 import seaborn as sn
3 plt.figure(figsize=(8,8))
4 sn.heatmap(confusion_matrix(y_actual_test, y_pred_test),annot=True,fmt='d')
5 plt.xlabel('Predicted Label')
6 plt.ylabel('Trut Label')
7 plt.title('Confusion Matrix of Sigmoid Test Data')

```

Out[52]: Text(0.5, 1.0, 'Confusion Matrix of Sigmoid Test Data')



```
In [51]: 1 c = confusion_matrix(y_actual_test, y_pred_test)
          2 print(c)
```

```
[ [ 963    0    3    1    0    1   11    1    0    0]
  [    0 1114   13    6    0    0    1    1    0    0]
  [   27   30  818   49   16    2   83    4    2    1]
  [   27   22   76  835    0   13    2   21   10    4]
  [    4   18    6    0  818    0   71    5    0   60]
  [  312    4   92  179    8  172   78   32    9    6]
  [   35    3   29    0    9    0  882    0    0    0]
  [   19   85   11   12   12    0    0  853    7   29]
  [   47   66  226  114   26   35   82   22  343   13]
  [   29   18    1   11  309    5   10  132    2  492]]
```

Observations:


- For $\eta = 0.01$, batch size = 64 and epochs = 15
 - Train Accuracy = 71%
 - Test Accuracy = 72.9%
- For $\eta = 0.04$, batch size = 64 and epochs = 15
 - Train Accuracy = 85.33%
 - Test Accuracy = 85.77%
- For $\eta = 0.01$, batch size = 64 and epochs = 10
 - Train Accuracy = 89.33%
 - Test Accuracy = 88.77%

TanH activation function

```
In [71]: 1 parameters_tanh, accuracy = NeuralNet(train_set, forward = 'TANH', learning_rate = 0.01, n_epochs = 10, batch_size = 100)
```

Epoch :- 1

Cost after iteration 0: 3.2504734281645433



```
In [72]: 1 print ("On the TEST set:")
2 y_actual_test, y_pred_test = Accuracy(test_set, parameters_tanh, 'TANH', 10000)
```

On the TEST set:

```
/home/mansi/.conda/envs/hbp/lib/python3.7/site-packages/ipykernel_launcher.py:11: RuntimeWarning: overflow encountered in exp
```

```
# This is added back by InteractiveShellApp.init_path()
```

```
/home/mansi/.conda/envs/hbp/lib/python3.7/site-packages/ipykernel_launcher.py:11: RuntimeWarning: invalid value encountered in true_divide
```

```
# This is added back by InteractiveShellApp.init_path()
```

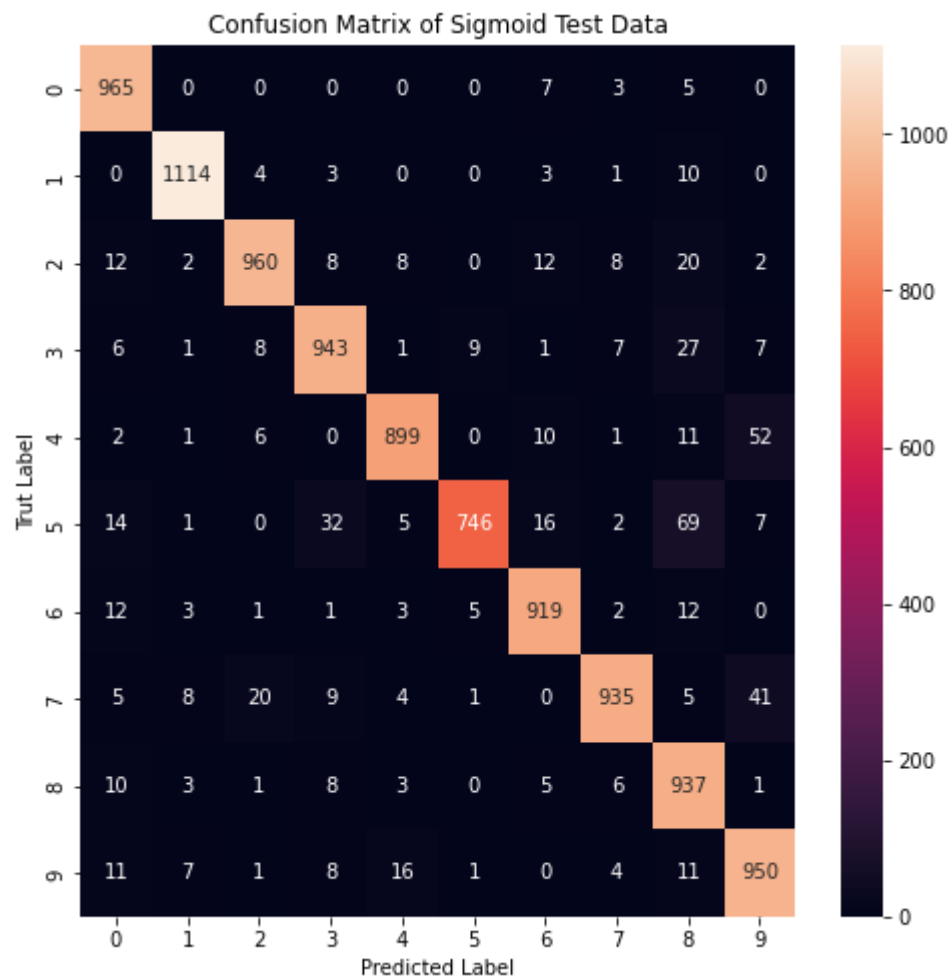
accuracy is =0.9368

```

In [73]: 1 #to get the heatmap for the confusion matrix
          2 import seaborn as sn
          3 plt.figure(figsize=(8,8))
          4 sn.heatmap(confusion_matrix(y_actual_test, y_pred_test),annot=True,fmt='d')
          5 plt.xlabel('Predicted Label')
          6 plt.ylabel('Trut Label')
          7 plt.title('Confusion Matrix of Sigmoid Test Data')

```

Out[73]: Text(0.5, 1.0, 'Confusion Matrix of Sigmoid Test Data')



```
In [74]: 1 c = confusion_matrix(y_actual_test, y_pred_test)
          2 print(c)
```

```
[ [ 965    0    0    0    0    0    7    3    5    0]
   [   0 1114    4    3    0    0    3    1   10    0]
   [  12    2   960    8    8    0   12    8   20    2]
   [   6    1    8   943    1    9    1    7   27    7]
   [   2    1    6    0   899    0   10    1   11   52]
   [  14    1    0   32    5   746   16    2   69    7]
   [  12    3    1    1    3    5   919    2   12    0]
   [   5    8   20    9    4    1    0   935    5   41]
   [  10    3    1    8    3    0    5    6   937    1]
   [  11    7    1    8   16    1    0    4   11   950]]
```

Observations:

- For $\eta = 0.01$, batch size = 64 and epochs = 15
 - Train Accuracy = 94.39%
 - Test Accuracy = 93.92%
- For $\eta = 0.04$, batch size = 64 and epochs = 15
 - Train Accuracy = 85.66%
 - Test Accuracy = 85.25%
- For $\eta = 0.01$, batch size = 64 and epochs = 10
 - Train Accuracy = 93.7%
 - Test Accuracy = 93.68%

Overall Observations:

- Model gives good accuracy in all the three activation functions with their respective best parameters. But the Pytorch model was much faster and still gives better accuracy than all the three here.
- ReLU and TanH give better accuracy as compared to Sigmoid

- After applying regularisation we donot observe much change in accuracy, but overall performance becomes better
- When we observe PyTorch Model,
 - we can tell that tanh doesnot perform as good as sigmoid and relu
 - but in our model tanh gives good results, probably tanh works good with gradient descent than with adam optimiser
 - ReLu shows the best results
 - overall the accuracy is high
- Uniform random Initialisation gave better results than random initialisation or zero initialisation