

```
In [ ]: 1 import numpy as np
        2 import matplotlib.pyplot as plt
        3 import torch
        4 import torchvision.datasets as data
        5 from torchvision.transforms import ToTensor
        6 from torch.utils.data import DataLoader
        7 import torch.nn as nn
        8 import torch.nn.functional as F
        9
       10 %matplotlib inline
       11 plt.rcParams['figure.figsize'] = (7.0, 4.0) # set default size of plots
```

```
In [ ]: 1 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
        2 device
```


Out[3]: device(type='cpu')

```
In [ ]: 1 train_set = data.MNIST(root = 'MNIST/raw/train-images-idx3-ubyte', train = True, transform= ToTensor(), down
        2 test_set = data.MNIST(root= 'MNIST/raw/train-images-idx3-ubyte', train= False, transform= ToTensor())
        3
        4 train_features = train_set.data
        5 train_labels = train_set.targets
```

```
In [ ]: 1
        2 # Display image and label.
        3
        4 print(f"Feature batch shape: {train_features.size()}")
        5 print(f"Labels batch shape: {train_labels.size()}")
        6 img = train_features[10].squeeze()
        7 label = train_labels[10]
        8 plt.imshow(img, cmap="gray")
        9 plt.show()
       10 print(f"Label: {label}")
```

Feature batch shape: torch.Size([60000, 28, 28])

Labels batch shape: torch.Size([60000])



Label: 3

```
In [ ]: 1 n_epochs = 5
        2 batch_size_train = 64
        3 batch_size_test = 1000
        4 learning_rate = 0.01
        5 input_size = (train_features.reshape(train_features.shape[0], -1)).shape[1]
        6 hidden_layer_1 = 500
        7 hidden_layer_2 = 250
        8 hidden_layer_3 = 100
        9 output_layer = 10
        10 momentum = 0.5
        11 log_interval = 10
        12
        13 random_seed = 1
        14 torch.backends.cudnn.enabled = False
        15 torch.manual_seed(random_seed)
```

Out[6]: <torch._C.Generator at 0x7ff6e8d6d510>

```
In [ ]: 1 train_loader = DataLoader(data.MNIST('/files/', train=True, download=True, transform = ToTensor()),batch_size=1)
        2
        3 test_loader = DataLoader(data.MNIST('/files/', train=False, download=True,transform=ToTensor()),batch_size=1)
```

```
In [ ]: 1 class classificationmodel_sigmoid(nn.Module):
2     def __init__(self, input_size, hidden_layer_1, hidden_layer_2, hidden_layer_3, output_layer):
3         super(classificationmodel_sigmoid,self).__init__()
4         self.linear1 = nn.Linear(input_size, hidden_layer_1)
5         self.linear2 = nn.Linear(hidden_layer_1, hidden_layer_2)
6         self.linear3 = nn.Linear(hidden_layer_2, hidden_layer_3)
7         self.linear4 = nn.Linear(hidden_layer_3, output_layer)
8         self.sigmoid = nn.Sigmoid()
9         self.softmax = nn.LogSoftmax(dim = 1)
10
11     def forward(self, image):
12         a = image.view(-1, input_size)
13         a = self.linear1(a)
14         a = self.sigmoid(a)
15         a = self.linear2(a)
16         a = self.sigmoid(a)
17         a = self.linear3(a)
18         a = self.sigmoid(a)
19         a = self.linear4(a)
20         a = self.softmax(a)
21         return a
```

```
In [ ]: 1 model_sigmoid = classificationmodel_sigmoid(input_size, hidden_layer_1, hidden_layer_2, hidden_layer_3, outp
```

```
In [ ]: 1 class classificationmodel_relu(nn.Module):
2     def __init__(self, input_size, hidden_layer_1, hidden_layer_2, hidden_layer_3, output_layer):
3         super(classificationmodel_relu,self).__init__()
4         self.linear1 = nn.Linear(input_size, hidden_layer_1)
5         self.linear2 = nn.Linear(hidden_layer_1, hidden_layer_2)
6         self.linear3 = nn.Linear(hidden_layer_2, hidden_layer_3)
7         self.linear4 = nn.Linear(hidden_layer_3, output_layer)
8         self.relu = nn.ReLU()
9         self.softmax = nn.LogSoftmax(dim = 1)
10
11     def forward(self, image):
12         a = image.view(-1, input_size)
13         a = self.linear1(a)
14         a = self.relu(a)
15         a = self.linear2(a)
16         a = self.relu(a)
17         a = self.linear3(a)
18         a = self.relu(a)
19         a = self.linear4(a)
20         a = self.softmax(a)
21         return a
```

```
In [ ]: 1 model_relu = classificationmodel_relu(input_size, hidden_layer_1, hidden_layer_2, hidden_layer_3, output_la
```

```
In [ ]: 1 class classificationmodel_tanh(nn.Module):
2     def __init__(self, input_size, hidden_layer_1, hidden_layer_2, hidden_layer_3, output_layer):
3         super(classificationmodel_tanh,self).__init__()
4         self.linear1 = nn.Linear(input_size, hidden_layer_1)
5         self.linear2 = nn.Linear(hidden_layer_1, hidden_layer_2)
6         self.linear3 = nn.Linear(hidden_layer_2, hidden_layer_3)
7         self.linear4 = nn.Linear(hidden_layer_3, output_layer)
8         self.tanh = nn.Tanh()
9         self.softmax = nn.LogSoftmax(dim = 1)
10
11     def forward(self, image):
12         a = image.view(-1, input_size)
13         a = self.linear1(a)
14         a = self.tanh(a)
15         a = self.linear2(a)
16         a = self.tanh(a)
17         a = self.linear3(a)
18         a = self.tanh(a)
19         a = self.linear4(a)
20         a = self.softmax(a)
21         return a
```

```
In [ ]: 1 model_tanh = classificationmodel_tanh(input_size, hidden_layer_1, hidden_layer_2, hidden_layer_3, output_la
```

```
In [ ]: 1 #Training
2 def train(model, optimizer):
3     n_steps = len(train_loader)
4     for e in range(n_epochs):
5         running_loss = 0
6         iter = 0
7         print(f'epoch = {e}')
8         for x, (images, labels) in enumerate(train_loader):
9             # Flatten the Image from 28*28 to 784 column vector
10            images = images.view(images.shape[0], -1)
11
12            # setting gradient to zeros
13            output = model(images)
14            loss = criterion(output, labels)
15            optimizer.zero_grad()
16            # backward propagation
17            loss.backward()
18            # update the gradient to new gradients
19            optimizer.step()
20            if (x+1)%100 == 0:
21                running_loss += loss.item()
22                iter +=1
23            if iter % 500 == 0:
24                # Calculate Accuracy
25                correct = 0
26                total = 0
27                # Iterate through test dataset
28                for images, labels in test_loader:
29                    # Load images to a Torch Variable
30                    images = images.view(images.shape[0], -1)
31
32                    # Forward pass only to get logits/output
33                    outputs = model(images)
34
35                    # Get predictions from the maximum value
36                    _, predicted = torch.max(outputs.data, 1)
37
38                    # Total number of labels
39                    total += labels.size(0)
40
41                    # Total correct predictions
```

```

42         correct += (predicted == labels).sum()
43
44         accuracy = 100 * correct / total
45
46         # Print Loss
47         print('Iteration: {}. Loss: {}. Accuracy: {}'.format(iter, loss.item(), accuracy))
48     print(f'epochs [{e+1}/{n_epochs}], Step[{x+1}/{n_steps}], Losses: {loss.item():.4f}')

```

```

In [ ]: 1 def accuracy(model, data_loader):
2         model.eval()
3         loss = 0
4         correct = 0
5         with torch.no_grad():
6             for image, label in data_loader:
7                 output = model(image)
8                 loss += criterion(output, label).item()
9                 pred = output.argmax(dim = 1, keepdim = True)
10                correct += pred.eq(label.view_as(pred)).sum().item()
11
12            loss /= len(data_loader.dataset)
13
14            print('\nAverage loss: {:.4f}, Accuracy: {}/{} ({:.0f}%)'.format(
15                loss, correct, len(data_loader.dataset),
16                100. * correct / len(data_loader.dataset)))
17

```

```

In [ ]: 1 criterion = nn.CrossEntropyLoss()
2         optimizer_sigmoid = torch.optim.Adam(model_sigmoid.parameters(), lr = learning_rate)
3         optimizer_relu = torch.optim.Adam(model_relu.parameters(), lr = learning_rate)
4         optimizer_tanh = torch.optim.Adam(model_tanh.parameters(), lr = learning_rate)

```

Training for Sigmoid Activation


```
In [ ]: 1 train(model_sigmoid, optimizer_sigmoid)

epoch = 0
Iteration: 500. Loss: 0.02223104238510132. Accuracy: 96.38999938964844
epoch = 1
Iteration: 500. Loss: 0.07424124330282211. Accuracy: 96.88999938964844
epoch = 2
Iteration: 500. Loss: 0.13215449452400208. Accuracy: 96.62000274658203
epoch = 3
Iteration: 500. Loss: 0.020750051364302635. Accuracy: 96.68000030517578
epoch = 4
Iteration: 500. Loss: 0.009185138158500195. Accuracy: 96.94000244140625
epochs [5/5], Step[938/938], Losses: 0.1411
```

```
In [ ]: 1 print('On the Train Set: ')
2 accuracy(model_sigmoid, train_loader)
3 print('On the Test Set: ')
4 accuracy(model_sigmoid, test_loader)
```

On the Train Set:

Average loss: 0.0010, Accuracy: 58813/60000 (98%)

On the Test Set:

Average loss: 0.0001, Accuracy: 9679/10000 (97%)

Training for ReLU activation

```
In [ ]: 1 train(model_relu, optimizer_relu)

epoch = 0
Iteration: 500. Loss: 0.16824673116207123. Accuracy: 93.94000244140625
epoch = 1
Iteration: 500. Loss: 0.08231014013290405. Accuracy: 96.4000015258789
epoch = 2
Iteration: 500. Loss: 0.2521880269050598. Accuracy: 96.16999816894531
epoch = 3
Iteration: 500. Loss: 0.20854493975639343. Accuracy: 96.41000366210938
epoch = 4
Iteration: 500. Loss: 0.04136047512292862. Accuracy: 96.1500015258789
epochs [5/5], Step[938/938], Losses: 0.1098
```

```
In [ ]: 1 print('On the Train Set: ')
2 accuracy(model_relu, train_loader)
3 print('On the Test Set: ')
4 accuracy(model_relu, test_loader)
```

On the Train Set:

Average loss: 0.0016, Accuracy: 58498/60000 (97%)

On the Test Set:

Average loss: 0.0002, Accuracy: 9620/10000 (96%)

Training for TanH Activation

```
In [ ]: 1 train(model_tanh, optimizer_tanh)

epoch = 0
Iteration: 500. Loss: 0.41016891598701477. Accuracy: 90.2300033569336
epoch = 1
Iteration: 500. Loss: 0.46348047256469727. Accuracy: 85.3499984741211
epoch = 2
Iteration: 500. Loss: 0.2446739375591278. Accuracy: 85.9800033569336
epoch = 3
Iteration: 500. Loss: 0.48340320587158203. Accuracy: 87.58999633789062
epoch = 4
Iteration: 500. Loss: 0.9405316114425659. Accuracy: 82.93000030517578
epochs [5/5], Step[938/938], Losses: 0.2936
```

```
In [ ]: 1 print('On the Train Set: ')
2 accuracy(model_tanh, train_loader)
3 print('On the Test Set: ')
4 accuracy(model_tanh, test_loader)
```

On the Train Set:


Average loss: 0.0069, Accuracy: 52729/60000 (88%)

On the Test Set:

Average loss: 0.0004, Accuracy: 8785/10000 (88%)

```
In [ ]: 1 def view_classify(img, ps):
2         ps = ps.data.numpy().squeeze()
3         fig, (ax1, ax2) = plt.subplots(figsize=(6,9), ncols=2)
4         ax1.imshow(img.resize_(1, 28, 28).numpy().squeeze())
5         ax1.axis('off')
6         ax2.barh(np.arange(10), ps)
7         ax2.set_aspect(0.1)
8         ax2.set_yticks(np.arange(10))
9         ax2.set_yticklabels(np.arange(10))
10        ax2.set_title('Class Probability')
11        ax2.set_xlim(0, 1.1)
12        plt.tight_layout()
```

```
In [ ]: 1 # Getting the image to test
2 images, labels = next(iter(train_loader))
3 # Flatten the image to pass in the model
4 img = images[0].view(1, 784)
5 # Turn off gradients to speed up this part
6 with torch.no_grad():
7     logps = model(img)
8 # Output of the network are log-probabilities, need to take exponential for probabilities
9 ps = torch.exp(logps)
10 view_classify(img, ps)
```



Observations:

Sigmoid gives accuracy as: ¶

- Training: 98%
- Testing: 97%

ReLU gives accuracy as:

- Training: 97%
- Testing: 96%

TanH gives accuracy as:

- Training: 88%
- Testing: 88%