

DEEP LEARNING MODEL ON MNIST DATASET

The architecture of this model comprises of three hidden layer with 500, 250 and 100 being their respective sizes

The aim is to model a complete handwritten digit recognizer, train it on the train set comprising 60,000 images and test performance on the test set comprising 10,000 images

Baseline model includes:

- Training data size of MNIST data = 60,000
 - with its respective batch size being 64
- Testing data size of MNIST data = 10,000
- Learning rate = 0.01
- number of epochs = 15

Libraries


```
In [4]: 1 import numpy as np
        2 import matplotlib.pyplot as plt
        3 import torch
        4 import torchvision.datasets as data
        5 from torchvision.transforms import ToTensor
        6 from torch.utils.data import DataLoader
        7 from sklearn.metrics import confusion_matrix, classification_report
        8
        9 %matplotlib inline
       10 plt.rcParams['figure.figsize'] = (7.0, 4.0) # set default size of plots
       11
       12 #Setting up GPU
       13 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
       14 device
```

```
Out[4]: device(type='cuda', index=0)
```

```
In [5]: 1 train_set = data.MNIST(root = 'MNIST/raw/train-images-idx3-ubyte', train = True, transform= ToTensor(), downlo
2 test_set = data.MNIST(root= 'MNIST/raw/train-images-idx3-ubyte', train= False, transform= ToTensor(), downlo
3
4 train_features = train_set.data
5 train_labels = train_set.targets
6 # Visualising data
7 # Display image and label
8
9 print(f"Feature batch shape: {train_features.size()}")
10 print(f"Labels batch shape: {train_labels.size()}")
11
12 fig = plt.figure()
13 for i in range(6):
14     plt.subplot(2,3,i+1)
15     plt.tight_layout()
16     plt.imshow(train_features[i].squeeze(), cmap='gray', interpolation='none')
17     plt.title("Ground Truth: {}".format(train_labels[i]))
18     plt.xticks([])
19     plt.yticks([])
20 plt.show()
21
```

Feature batch shape: torch.Size([60000, 28, 28])

Labels batch shape: torch.Size([60000])



```
In [6]: 1 batchsize = 64
        2 input_size = (train_features.reshape(train_features.shape[0], -1)).shape[1]
        3 hidden_layer_1 = 500
        4 hidden_layer_2 = 250
        5 hidden_layer_3 = 100
        6 output_layer = 10
        7 learning_rate = 0.01
```

```

In [34]: 1 def one_hot_encode(Y):
2     output = np.eye(10)[np.array(Y).reshape(-1)]
3     return output.reshape(list(np.shape(Y))+[10])
4
5
6 def tanh(x):
7     return ((np.exp(x)-np.exp(-x))/(np.exp(x)+np.exp(-x)))
8
9
10 def softmax(x):
11     return np.exp(x)/sum(np.exp(x))
12
13 def derivative_tanh(x):
14     t = activation(x, 'TANH')
15     return (1-t**2)
16
17 def data_flattening(features, labels, one_hot = True):
18     features = features.numpy()
19     labels = labels.numpy()
20     X = (features.reshape(features.shape[0], -1))
21     if one_hot:
22         Y = one_hot_encode(labels)
23     else:
24         Y = labels
25     return X, Y
26
27 #Function to Initialise the parameters
28
29 def initialise_parameter(dim):
30     np.random.seed(11)
31
32     parameters = {}
33     L = len(dim)
34     for i in range(1, L):
35         Ni = dim[i-1]
36         No = dim[i]
37         M = np.sqrt(6/(Ni+No))
38         parameters["W" + str(i)] = np.asarray(np.random.uniform(-M, M,size = (No,Ni)))
39         parameters["b" + str(i)] = np.zeros((dim[i], 1))
40
41     assert(parameters["W" + str(i)].shape == (dim[i], dim[i-1]))

```

```

42     assert(parameters["b" + str(i)].shape == (dim[i], 1))
43     return parameters
44
45
46 def forward_propagation(X,parameters):
47     W1=parameters["W1"]
48     b1=parameters["b1"]
49     W2=parameters["W2"]
50     b2=parameters["b2"]
51     W3=parameters["W3"]
52     b3=parameters["b3"]
53     W4=parameters["W4"]
54     b4=parameters["b4"]
55     Z1=np.dot(W1,X.T)+b1
56     A1=tanh(Z1)
57     Z2 = np.dot(W2,A1) + b2
58     A2=tanh(Z2)
59     Z3 = np.dot(W3,A2) + b3
60     A3=tanh(Z3)
61     Z4=np.dot(W4,A3)+b4
62     A4=softmax(Z4)
63
64     cache = (Z1, A1, W1, b1, Z2, A2, W2, b2, Z3, A3, W3, b3 , Z4 ,A4 , W4 , b4)
65
66     return A4,cache
67
68 # GRADED FUNCTION: backward_propagation_with_regularization
69
70 def backward_propagation(X, Y, cache, lambd):
71
72     m = batchsize
73     (Z1, A1, W1, b1, Z2, A2, W2, b2, Z3, A3, W3, b3 , Z4 ,A4 , W4 , b4) = cache
74
75     dZ4 = A4 - Y
76     dW4 = (1./m)*((np.dot(dZ4, A3.T)))+((lambd/m)*W4)
77     db4 = 1./m * np.sum(dZ4, axis=1, keepdims = True)
78
79     dA3 = np.dot(W4.T, dZ4)
80     dZ3 = np.multiply(dA3, derivative_tanh(A3))
81     dW3 = (1./m)*((np.dot(dZ3, A2.T)))+((lambd/m)*W3)
82     db3 = 1./m * np.sum(dZ3, axis=1, keepdims = True)
83

```

```

84     dA2 = np.dot(W3.T, dZ3)
85     dZ2 = np.multiply(dA2, derivative_tanh(A2))
86     dW2 = (1./m)*((np.dot(dZ2, A1.T)))+(lambda/m)*W2)
87     db2 = 1./m * np.sum(dZ2, axis=1, keepdims = True)
88
89     dA1 = np.dot(W2.T, dZ2)
90     dZ1 = np.multiply(dA1, derivative_tanh(A1))
91     dW1 = (1./m)*((np.dot(dZ1, X)))+(lambda/m)*W1)
92     db1 = 1./m * np.sum(dZ1, axis=1, keepdims = True)
93
94     gradients = {"dZ4": dZ4, "dW4": dW4, "db4": db4,
95                  "dA3": dA3, "dZ3": dZ3, "dW3": dW3, "db3": db3,
96                  "dA2": dA2, "dZ2": dZ2, "dW2": dW2, "db2": db2,
97                  "dA1": dA1, "dZ1": dZ1, "dW1": dW1, "db1": db1}
98
99     return gradients
100
101 def compute_cost(A, Y,cache, lambda):
102     #A is predicted
103     #Y is actual
104     m = Y.shape[1]
105     (Z1, A1, W1, b1, Z2, A2, W2, b2, Z3, A3, W3, b3 , Z4 ,A4 , W4 , b4) = cache
106     logprobs = np.multiply(-np.log(A),Y) + np.multiply(-np.log(1 - A), 1 - Y)
107     cost = 1./m * np.nansum(logprobs)
108     L2_regularization_cost = (lambda/(2*m))*(np.sum(np.square(W1))+np.sum(np.square(W2))+np.sum(np.square(W3)
109     costTotal = cost + L2_regularization_cost
110     return costTotal
111
112 #Function to update the parameters
113
114 def update_parameters(parameters, grads, learning_rate):
115     W1 = parameters["W1"]
116     b1 = parameters["b1"]
117     W2 = parameters["W2"]
118     b2 = parameters["b2"]
119     W3 = parameters["W3"]
120     b3 = parameters["b3"]
121     W4 = parameters["W4"]
122     b4 = parameters["b4"]
123     dW1 = grads["dW1"]
124     db1 = grads["db1"]
125     dW2 = grads["dW2"]

```

```
126     db2 = grads["db2"]
127     dW3 = grads["dW3"]
128     db3 = grads["db3"]
129     dW4 = grads["dW4"]
130     db4 = grads["db4"]
131     W1 = W1 - learning_rate * dW1
132     b1 = b1 - learning_rate * db1
133     W2 = W2 - learning_rate * dW2
134     b2 = b2 - learning_rate * db2
135     W3 = W3 - learning_rate * dW3
136     b3 = b3 - learning_rate * db3
137     W4 = W4 - learning_rate * dW4
138     b4 = b4 - learning_rate * db4
139
140     parameters={"W1":W1, "b1":b1,
141                "W2":W2, "b2":b2,
142                "W3":W3, "b3":b3,
143                "W4":W4, "b4":b4}
144     return parameters
145
146 # Predict Labels
147
148 def Accuracy(dataset, parameters, size):
149
150     features = dataset.data
151     labels = dataset.targets
152     X, Y = data_flattening(features, labels, one_hot = False)
153     y = Y.T
154     p = np.zeros(size, dtype = int)
155     # Forward propagation
156     a4, caches = forward_propagation(X, parameters)
157     p = np.argmax(a4, axis = 0)
158     a = np.mean((p == y))
159
160     print("accuracy is =" + str(a))
161
162     return y, p
163
164 def find_accuracy(y_actual, y_pred):
165     accuracy = np.count_nonzero(np.argmax(y_pred, axis=0) == np.argmax(y_actual, axis=1)) / y_actual.shape[0]
166     return accuracy
167
```



```
168 def predict(X,Y,parameters):
169     """
170     This function is used to predict the results of a n-layer neural network.
171
172     Arguments:
173     X -- data set of examples you would like to label
174     Y -- data set of examples
175     parameters -- parameters of the trained model
176
177     Returns:
178     ypred -- predictions for the given dataset X
179     """
180
181
182     y_pred,cache=forward_propagation(X,parameters)
183     return y_pred
```

```

In [35]: 1 def model(dataset, numEpochs, lambd):
2
3     k=len(dataset)                                #length of the dataset
4     numBatches=k/batchsize
5     layers_dims = [input_size, hidden_layer_1, hidden_layer_2, hidden_layer_3, output_layer]# number of bat
6     parameters=initialise_parameter(layers_dims)    # initializing the parameters
7     costs=[]                                       #cost accumulation
8     acc=[]                                         #accuracy
9     for epoch in range(numEpochs):
10         for j in range(int(numBatches)):
11             # Data loader
12             loader = DataLoader(dataset=dataset,batch_size = batchsize ,shuffle=True)
13             dataiter = iter(loader)
14             data = next(dataiter)
15             features, labels = data
16
17             X, Y = data_flattening(features, labels, True)
18             y_pred,cache=forward_propagation(X,parameters)
19             cost=compute_cost(y_pred,Y.T, cache, lambd)
20             gradients=backward_propagation(X,Y.T,cache, lambd)
21             parameters=update_parameters(parameters,gradients,learning_rate)
22             if j%200 ==0:
23                 print (f'Epoch [{epoch+1}/{numEpochs}], Step [{j+1}/{int(numBatches)}], Loss: {cost.item():
24                 acc.append(find_accuracy(Y,y_pred))
25                 costs.append(cost)
26     return parameters, costs, acc

```

```
In [41]: 1 train_parameters, train_costs, train_acc = model(train_set,15, 0.7)
```

```
Epoch [1/15], Step [1/937], Loss: 9.2180
Epoch [1/15], Step [201/937], Loss: 6.9773
Epoch [1/15], Step [401/937], Loss: 6.6116
Epoch [1/15], Step [601/937], Loss: 6.2354
Epoch [1/15], Step [801/937], Loss: 5.8071
Epoch [2/15], Step [1/937], Loss: 5.9277
Epoch [2/15], Step [201/937], Loss: 5.3748
Epoch [2/15], Step [401/937], Loss: 5.1373
Epoch [2/15], Step [601/937], Loss: 5.0750
Epoch [2/15], Step [801/937], Loss: 4.8159
Epoch [3/15], Step [1/937], Loss: 4.6255
Epoch [3/15], Step [201/937], Loss: 4.6047
Epoch [3/15], Step [401/937], Loss: 4.5472
Epoch [3/15], Step [601/937], Loss: 4.4994
Epoch [3/15], Step [801/937], Loss: 4.1898
Epoch [4/15], Step [1/937], Loss: 3.9538
Epoch [4/15], Step [201/937], Loss: 3.8346
Epoch [4/15], Step [401/937], Loss: 4.0730
Epoch [4/15], Step [601/937], Loss: 3.4899
Epoch [4/15], Step [801/937], Loss: 3.6199
Epoch [5/15], Step [1/937], Loss: 3.4327
Epoch [5/15], Step [201/937], Loss: 3.2664
Epoch [5/15], Step [401/937], Loss: 3.1833
Epoch [5/15], Step [601/937], Loss: 3.0535
Epoch [5/15], Step [801/937], Loss: 2.8444
Epoch [6/15], Step [1/937], Loss: 3.1384
Epoch [6/15], Step [201/937], Loss: 2.6813
Epoch [6/15], Step [401/937], Loss: 2.8368
Epoch [6/15], Step [601/937], Loss: 2.5131
Epoch [6/15], Step [801/937], Loss: 2.3151
Epoch [7/15], Step [1/937], Loss: 2.6430
Epoch [7/15], Step [201/937], Loss: 2.2990
Epoch [7/15], Step [401/937], Loss: 2.3621
Epoch [7/15], Step [601/937], Loss: 2.2973
Epoch [7/15], Step [801/937], Loss: 2.4129
Epoch [8/15], Step [1/937], Loss: 2.4792
Epoch [8/15], Step [201/937], Loss: 2.0638
Epoch [8/15], Step [401/937], Loss: 2.0477
Epoch [8/15], Step [601/937], Loss: 2.0168
```

```
Epoch [8/15], Step [801/937], Loss: 2.0216
Epoch [9/15], Step [1/937], Loss: 1.9025
Epoch [9/15], Step [201/937], Loss: 1.9421
Epoch [9/15], Step [401/937], Loss: 1.6988
Epoch [9/15], Step [601/937], Loss: 2.0338
Epoch [9/15], Step [801/937], Loss: 2.0210
Epoch [10/15], Step [1/937], Loss: 1.9000
Epoch [10/15], Step [201/937], Loss: 1.7415
Epoch [10/15], Step [401/937], Loss: 1.8228
Epoch [10/15], Step [601/937], Loss: 1.7201
Epoch [10/15], Step [801/937], Loss: 1.4069
Epoch [11/15], Step [1/937], Loss: 1.5675
Epoch [11/15], Step [201/937], Loss: 1.7428
Epoch [11/15], Step [401/937], Loss: 1.6723
Epoch [11/15], Step [601/937], Loss: 1.6761
Epoch [11/15], Step [801/937], Loss: 1.4715
Epoch [12/15], Step [1/937], Loss: 1.7129
Epoch [12/15], Step [201/937], Loss: 1.4249
Epoch [12/15], Step [401/937], Loss: 1.4464
Epoch [12/15], Step [601/937], Loss: 1.4510
Epoch [12/15], Step [801/937], Loss: 1.2633
Epoch [13/15], Step [1/937], Loss: 1.2307
Epoch [13/15], Step [201/937], Loss: 1.2086
Epoch [13/15], Step [401/937], Loss: 1.3936
Epoch [13/15], Step [601/937], Loss: 1.1806
Epoch [13/15], Step [801/937], Loss: 1.2469
Epoch [14/15], Step [1/937], Loss: 1.1430
Epoch [14/15], Step [201/937], Loss: 1.4050
Epoch [14/15], Step [401/937], Loss: 1.2918
Epoch [14/15], Step [601/937], Loss: 1.3070
Epoch [14/15], Step [801/937], Loss: 1.1565
Epoch [15/15], Step [1/937], Loss: 1.1579
Epoch [15/15], Step [201/937], Loss: 1.4134
Epoch [15/15], Step [401/937], Loss: 1.0654
Epoch [15/15], Step [601/937], Loss: 1.0153
Epoch [15/15], Step [801/937], Loss: 1.1690
```

```
In [43]: 1 train_acc[-1]
```

```
Out[43]: 0.953125
```

```
In [44]: 1 print ("On the TEST set:")  
        2 y_actual_test, y_pred_test = Accuracy(test_set, trained_parameters, 10000)
```


```
On the TEST set:  
accuracy is =0.9095
```

```

In [45]: 1 #to get the heatmap for the confusion matrix
2 import seaborn as sn
3 plt.figure(figsize=(8,8))
4 sn.heatmap(confusion_matrix(y_actual_test, y_pred_test),annot=True,fmt='d')
5 plt.xlabel('Predicted Label')
6 plt.ylabel('Trut Label')
7 plt.title('Confusion Matrix of Sigmoid Test Data')

```

Out[45]: Text(0.5, 1.0, 'Confusion Matrix of Sigmoid Test Data')



```
In [46]: 1 c = confusion_matrix(y_actual_test, y_pred_test)
          2 print(c)
```

```
[[ 964    0    0    1    0    0    6    1    8    0]
 [   0 1087    3    3    1    0    4    1   35    1]
 [   13    3  929    9    6    0   15    7   42    8]
 [    8    0   18  918    0   10    2    7   44    3]
 [    1    3    4    0  839    0   15    1   25   94]
 [   27    1    1   47    3  644   17    4  142    6]
 [   15    3    2    0    7    2  913    0   16    0]
 [    3   10   25    9    6    0    0  913    8   54]
 [    5    1    3    9    4    1    5    2  941    3]
 [    8    2    0    6    7    1    1    4   33  947]]
```

Observations:

- Tanh showed good accuracy :
 - for trainset - 94.39%
 - for testset - 93.92%
- After applying regularization on the trainset with $\lambda = 0.7$ and number of epochs = 15 we get accuracy as :
 - for trainset - 95.31%
 - for testset - 90.35%

```
In [ ]: 1
```