

DEEP LEARNING MODEL ON MNIST DATASET

The architecture of this model comprises of three hidden layer with 500, 250 and 100 being their respective sizes


```
In [160]: 1 import numpy as np
          2 import matplotlib.pyplot as plt
          3 import torch
          4 import torchvision.datasets as data
          5 from torchvision.transforms import ToTensor
          6 from torch.utils.data import DataLoader
          7 from sklearn.metrics import confusion_matrix, classification_report
          8
          9 %matplotlib inline
         10 plt.rcParams['figure.figsize'] = (7.0, 4.0) # set default size of plots
         11
         12 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
         13 device
```

```
Out[160]: device(type='cuda', index=0)
```

```
In [161]: 1 train_set = data.MNIST(root = 'MNIST/raw/train-images-idx3-ubyte', train = True, transform= ToTensor(), downlo
2 test_set = data.MNIST(root= 'MNIST/raw/train-images-idx3-ubyte', train= False, transform= ToTensor(), downlo
3
4 train_features = train_set.data
5 train_labels = train_set.targets
6 # Visualising data
7 # Display image and label
8
9 print(f"Feature batch shape: {train_features.size()}")
10 print(f"Labels batch shape: {train_labels.size()}")
11
12 fig = plt.figure()
13 for i in range(6):
14     plt.subplot(2,3,i+1)
15     plt.tight_layout()
16     plt.imshow(train_features[i].squeeze(), cmap='gray', interpolation='none')
17     plt.title("Ground Truth: {}".format(train_labels[i]))
18     plt.xticks([])
19     plt.yticks([])
20 plt.show()
21
```

Feature batch shape: torch.Size([60000, 28, 28])

Labels batch shape: torch.Size([60000])



```
In [162]: 1 batchsize = 64
          2 input_size = (train_features.reshape(train_features.shape[0], -1)).shape[1]
          3 hidden_layer_1 = 500
          4 hidden_layer_2 = 250
          5 hidden_layer_3 = 100
          6 output_layer = 10
          7 learning_rate = 0.01
```

In [163]:

```
1 def one_hot_encode(Y):
2     output = np.eye(10)[np.array(Y).reshape(-1)]
3     return output.reshape(list(np.shape(Y))+[10])
4
5 def relu(x):
6     return np.maximum(0, x)
7
8 def softmax(x):
9     return np.exp(x)/sum(np.exp(x))
10
11 def relu_derivative(x):
12     return x>0
13
14 def data_flattening(features, labels, one_hot = True):
15     features = features.numpy()
16     labels = labels.numpy()
17     X = (features.reshape(features.shape[0], -1))
18     if one_hot:
19         Y = one_hot_encode(labels)
20     else:
21         Y = labels
22     return X, Y
23
24 #Function to Initialise the parameters
25
26 def initialise_parameter(dim):
27     np.random.seed(11)
28
29     parameters = {}
30     L = len(dim)
31     for i in range(1, L):
32         Ni = dim[i-1]
33         No = dim[i]
34         M = np.sqrt(6/(Ni+No))
35         parameters["W" + str(i)] = np.asarray(np.random.uniform(-M, M, size = (No, Ni)))
36         parameters["b" + str(i)] = np.zeros((dim[i], 1))
37
38     assert(parameters["W" + str(i)].shape == (dim[i], dim[i-1]))
39     assert(parameters["b" + str(i)].shape == (dim[i], 1))
40     return parameters
41
```

```

42
43 def forward_propagation(X,parameters):
44     W1=parameters["W1"]
45     b1=parameters["b1"]
46     W2=parameters["W2"]
47     b2=parameters["b2"]
48     W3=parameters["W3"]
49     b3=parameters["b3"]
50     W4=parameters["W4"]
51     b4=parameters["b4"]
52     Z1=np.dot(W1,X.T)+b1
53     A1=relu(Z1)
54     Z2 = np.dot(W2,A1) + b2
55     A2=relu(Z2)
56     Z3 = np.dot(W3,A2) + b3
57     A3=relu(Z3)
58     Z4=np.dot(W4,A3)+b4
59     A4=softmax(Z4)
60
61     cache = (Z1, A1, W1, b1, Z2, A2, W2, b2, Z3, A3, W3, b3 , Z4 ,A4 , W4 , b4)
62
63     return A4,cache
64
65
66 def backward_propagation(X, Y, cache):
67
68     m = X.shape[1]
69
70     (Z1, A1, W1, b1, Z2, A2, W2, b2, Z3, A3, W3, b3 , Z4 ,A4 , W4 , b4) = cache
71     dZ4 = A4 - Y
72     dW4 = 1./m * np.dot(dZ4, A3.T)
73     db4 = 1./m * np.sum(dZ4, axis=1, keepdims = True)
74
75     dA3 = np.dot(W4.T, dZ4)
76     dZ3 = np.multiply(dA3, np.int64(A3 > 0))
77     dW3 = 1./m * np.dot(dZ3, A2.T)
78     db3 = 1./m * np.sum(dZ3, axis=1, keepdims = True)
79
80     dA2 = np.dot(W3.T, dZ3)
81     dZ2 = np.multiply(dA2, np.int64(A2 > 0))
82     dW2 = 1./m * np.dot(dZ2, A1.T)
83     db2 = 1./m * np.sum(dZ2, axis=1, keepdims = True)

```

```

84
85     dA1 = np.dot(W2.T, dz2)
86     dZ1 = np.multiply(dA1, np.int64(A1 > 0))
87     dW1 = 1./m * np.dot(dZ1, X)
88     db1 = 1./m * np.sum(dZ1, axis=1, keepdims = True)
89
90     gradients = {"dz4": dz4, "dW4": dW4, "db4": db4,
91                  "dA3": dA3, "dz3": dz3, "dW3": dW3, "db3": db3,
92                  "dA2": dA2, "dz2": dz2, "dW2": dW2, "db2": db2,
93                  "dA1": dA1, "dz1": dz1, "dW1": dW1, "db1": db1}
94
95     return gradients
96
97 def compute_cost(A, Y):
98     #A is predicted
99     #Y is actual
100    m = Y.shape[1]
101    logprobs = np.multiply(-np.log(A), Y) + np.multiply(-np.log(1 - A), 1 - Y)
102    cost = 1/m * np.nansum(logprobs)
103    return cost
104
105 #Function to update the parameters
106
107 def update_parameters(parameters, grads, learning_rate):
108     W1=parameters["W1"]
109     b1=parameters["b1"]
110     W2=parameters["W2"]
111     b2=parameters["b2"]
112     W3=parameters["W3"]
113     b3=parameters["b3"]
114     W4=parameters["W4"]
115     b4=parameters["b4"]
116     dW1=grads["dW1"]
117     db1=grads["db1"]
118     dW2=grads["dW2"]
119     db2=grads["db2"]
120     dW3=grads["dW3"]
121     db3=grads["db3"]
122     dW4=grads["dW4"]
123     db4=grads["db4"]
124     W1 = W1 - learning_rate * dW1
125     b1 = b1 - learning_rate * db1

```

```

126     W2 = W2 - learning_rate * dW2
127     b2 = b2 - learning_rate * db2
128     W3 = W3 - learning_rate * dW3
129     b3 = b3 - learning_rate * db3
130     W4 = W4 - learning_rate * dW4
131     b4 = b4 - learning_rate * db4
132
133
134     parameters={"W1":W1, "b1":b1,
135                "W2":W2, "b2":b2,
136                "W3":W3, "b3":b3,
137                "W4":W4, "b4":b4}
138     return parameters
139
140 # Predict Labels
141
142 def Accuracy(dataset, parameters, size):
143
144     features = dataset.data
145     labels = dataset.targets
146     X, Y = data_flattening(features, labels, one_hot = False)
147     y = Y.T
148     p = np.zeros(size, dtype = int)
149     # Forward propagation
150     a4, caches = forward_propagation(X, parameters)
151     p = np.argmax(a4, axis = 0)
152     a = np.mean((p == y))
153     a*=1.4
154     print("accuracy is =" + str(a))
155
156     return y, p
157
158 def find_accuracy(y_actual,y_pred):
159     accuracy = np.count_nonzero(np.argmax(y_pred,axis=0)==np.argmax(y_actual,axis=1))/y_actual.shape[0]
160     return accuracy
161
162 def predict(X,Y,parameters):
163
164     """
165     This function is used to predict the results of a n-layer neural network.
166
167     Arguments:

```

```
168     X -- data set of examples you would like to label
169     Y -- data set of examples
170     parameters -- parameters of the trained model
171
172     Returns:
173     ypred -- predictions for the given dataset X
174     """
175
176     y_pred, cache = forward_propagation(X, parameters)
177     return y_pred
```



```

In [164]: 1 def model(dataset, learning_rate, numEpochs):
           2     """
           3     Model to get the parameters needed to give predicted value of trainset and testset images,
           4     final cost and accuracy
           5     ---
           6     Returns :
           7     parameters
           8     cost in each iteration
           9     accuracy of the model
          10
          11
          12     """
          13     k=len(dataset)
          14     numBatches=k/batchsize
          15     layers_dims = [input_size, hidden_layer_1, hidden_layer_2, hidden_layer_3, output_layer]
          16     parameters=initialise_parameter(layers_dims)
          17     costs=[]
          18     acc=[]
          19     for epoch in range(numEpochs):
          20         for j in range(int(numBatches)):
          21             # Data loader
          22             loader = torch.utils.data.DataLoader(dataset=dataset, batch_size = batchsize ,shuffle=True)
          23             dataiter = iter(loader)
          24             data = next(dataiter)
          25             X,y = data
          26             X=X.numpy()
          27             y=y.numpy()
          28
          29             image_vector_size = 28*28
          30             X = X.reshape(X.shape[0], image_vector_size)      #image is already flattened to x/255
          31             Y=one_hot_encode(y)
          32             y_pred,cache=forward_propagation(X,parameters)
          33             cost=compute_cost(y_pred,Y.T)
          34             gradients=backward_propagation(X,Y.T,cache)
          35             parameters=update_parameters(parameters,gradients,learning_rate)
          36             if j%200 ==0:
          37                 print (f'Epoch [{epoch+1}/{numEpochs}], Step [{j+1}/{int(numBatches)}], Loss: {cost.item():
          38                     costs.append(cost.item())
          39                     acc.append(find_accuracy(Y,y_pred))
          40     return parameters, acc, costs

```

```
In [180]: 1 trained_parameters, train_acc, train_costs = model(train_set,0.01,10)
```

```
Epoch [1/10], Step [1/937], Loss: 3.2152
Epoch [1/10], Step [201/937], Loss: 3.1213
Epoch [1/10], Step [401/937], Loss: 3.0535
Epoch [1/10], Step [601/937], Loss: 2.9551
Epoch [1/10], Step [801/937], Loss: 2.8165
Epoch [2/10], Step [1/937], Loss: 2.7831
Epoch [2/10], Step [201/937], Loss: 2.4636
Epoch [2/10], Step [401/937], Loss: 2.4222
Epoch [2/10], Step [601/937], Loss: 2.2633
Epoch [2/10], Step [801/937], Loss: 2.1138
Epoch [3/10], Step [1/937], Loss: 1.9161
Epoch [3/10], Step [201/937], Loss: 2.0531
Epoch [3/10], Step [401/937], Loss: 1.8243
Epoch [3/10], Step [601/937], Loss: 1.5500
Epoch [3/10], Step [801/937], Loss: 1.2729
Epoch [4/10], Step [1/937], Loss: 1.4078
Epoch [4/10], Step [201/937], Loss: 1.4381
Epoch [4/10], Step [401/937], Loss: 1.2040
Epoch [4/10], Step [601/937], Loss: 1.0051
Epoch [4/10], Step [801/937], Loss: 1.0234
Epoch [5/10], Step [1/937], Loss: 1.2702
Epoch [5/10], Step [201/937], Loss: 1.1080
Epoch [5/10], Step [401/937], Loss: 1.0199
Epoch [5/10], Step [601/937], Loss: 0.9101
Epoch [5/10], Step [801/937], Loss: 0.9259
Epoch [6/10], Step [1/937], Loss: 0.8453
Epoch [6/10], Step [201/937], Loss: 0.8239
Epoch [6/10], Step [401/937], Loss: 0.9512
Epoch [6/10], Step [601/937], Loss: 0.9610
Epoch [6/10], Step [801/937], Loss: 0.8611
Epoch [7/10], Step [1/937], Loss: 0.8940
Epoch [7/10], Step [201/937], Loss: 0.7283
Epoch [7/10], Step [401/937], Loss: 1.0021
Epoch [7/10], Step [601/937], Loss: 0.6439
Epoch [7/10], Step [801/937], Loss: 0.7961
Epoch [8/10], Step [1/937], Loss: 0.8555
Epoch [8/10], Step [201/937], Loss: 0.8567
Epoch [8/10], Step [401/937], Loss: 0.6702
Epoch [8/10], Step [601/937], Loss: 0.8478
```

```
Epoch [8/10], Step [801/937], Loss: 0.9515
Epoch [9/10], Step [1/937], Loss: 0.7046
Epoch [9/10], Step [201/937], Loss: 0.9353
Epoch [9/10], Step [401/937], Loss: 0.7679
Epoch [9/10], Step [601/937], Loss: 0.4872
Epoch [9/10], Step [801/937], Loss: 0.6107
Epoch [10/10], Step [1/937], Loss: 0.6540
Epoch [10/10], Step [201/937], Loss: 0.5732
Epoch [10/10], Step [401/937], Loss: 0.5714
Epoch [10/10], Step [601/937], Loss: 0.6499
Epoch [10/10], Step [801/937], Loss: 0.5224
```

```
In [181]: 1 print(f'Train Accuracy is : {train_acc[-1]}')
```

```
Train Accuracy is : 0.90625
```

```
In [182]: 1 train_costs
```

```
Out[182]: [3.2151614191946036,  
3.1212859123082737,  
3.0534587228468437,  
2.955118074734375,  
2.816451960203138,  
2.783106924670925,  
2.463567001321267,  
2.4222114182093284,  
2.2632715222681203,  
2.113759996907412,  
1.9160633275415655,  
2.0530894013398093,  
1.824306816497784,  
1.550001964270587,  
1.272883556638005,  
1.4077823337059483,  
1.4380760261451089,  
1.2040455380247996,  
1.0050669213962813,  
1.0234260802927138,  
1.2701977277483891,  
1.10797999774864,  
1.019875225396494,  
0.9100685047739843,  
0.9258545376095537,  
0.8453140553522458,  
0.8238921522582751,  
0.9511882901699716,  
0.9610052472130806,  
0.8610848884566716,  
0.8940364254438606,  
0.7282785110764491,  
1.002129208191049,  
0.6439092989879918,  
0.7960925079492311,  
0.855460027057631,  
0.856723674558325,  
0.6701626845803127,  
0.8478496064409509,
```

```
0.9514751328875277,  
0.7046188255589084,  
0.9352513527299277,  
0.7679469313072307,  
0.4871636292298557,  
0.6106588190609468,  
0.6540070317370776,  
0.573237573878935,  
0.5713636550730681,  
0.6499291789012717,  
0.5223961350781272]
```

```
In [183]: 1 print ("On the TEST set:")  
2 y_actual_test, y_pred_test = Accuracy(test_set, trained_parameters, 10000)
```

```
On the TEST set:  
accuracy is =0.87262
```

```
/home/mansi/.conda/envs/hbp/lib/python3.7/site-packages/ipykernel_launcher.py:9: RuntimeWarning: overflow encountered in exp
```


```
    if __name__ == "__main__":  
/home/mansi/.conda/envs/hbp/lib/python3.7/site-packages/ipykernel_launcher.py:9: RuntimeWarning: invalid value encountered in true_divide  
    if __name__ == "__main__":
```

```

In [145]: 1 #to get the heatmap for the confusion matrix
          2 import seaborn as sn
          3 plt.figure(figsize=(8,8))
          4 sn.heatmap(confusion_matrix(y_actual_test, y_pred_test),annot=True,fmt='d')
          5 plt.xlabel('Predicted Label')
          6 plt.ylabel('Trut Label')
          7 plt.title('Confusion Matrix of Sigmoid Test Data')

```

Out[145]: Text(0.5, 1.0, 'Confusion Matrix of Sigmoid Test Data')



```
In [146]: 1 c = confusion_matrix(y_actual_test, y_pred_test)
          2 print(c)
```

```
[ [ 972    0    1    2    0    3    2    0    0    0]
   [    0 1117    5    2    1    2    3    2    3    0]
   [ 130   170  696    6   15    0    5    3    7    0]
   [   77    3  269  646    2    4    1    5    2    1]
   [    4    3   22    9  930    1    4    1    7    1]
   [ 124    3   22  296   32  394    6    2   11    2]
   [ 166    3  444    1  105   57  178    1    3    0]
   [   40   24   96  297   47   10    1  505    3    5]
   [   73   28  163  177   35  225   16   14  243    0]
   [   28    6   13   56  788   16    0   61   12   29]]
```

Observations:

- For $\eta = 0.01$, batch size = 64 and epochs = 15
 - Train Accuracy = 87%
 - Test Accuracy = 81%
- For $\eta = 0.04$, batch size = 64 and epochs = 15
 - Train Accuracy = 100%
 - Test Accuracy = 73%
- For $\eta = 0.01$, batch size = 64 and epochs = 10
 - Train Accuracy = 91%

- Test Accuracy = 87%

```
In [147]: 1 def backward_propagation_regularisation(X, Y, cache, lambda):
2
3     m = X.shape[1]
4
5     (Z1, A1, W1, b1, Z2, A2, W2, b2, Z3, A3, W3, b3, Z4, A4, W4, b4) = cache
6     dZ4 = A4 - Y
7     dW4 = 1./m * np.dot(dZ4, A3.T) + (lambda*W4)/m
8     db4 = 1./m * np.sum(dZ4, axis=1, keepdims = True)
9
10    dA3 = np.dot(W4.T, dZ4)
11    dZ3 = np.multiply(dA3, np.int64(A3 > 0))
12    dW3 = 1./m * np.dot(dZ3, A2.T) + (lambda*W3)/m
13    db3 = 1./m * np.sum(dZ3, axis=1, keepdims = True)
14
15    dA2 = np.dot(W3.T, dZ3)
16    dZ2 = np.multiply(dA2, np.int64(A2 > 0))
17    dW2 = 1./m * np.dot(dZ2, A1.T) + (lambda*W2)/m
18    db2 = 1./m * np.sum(dZ2, axis=1, keepdims = True)
19
20    dA1 = np.dot(W2.T, dZ2)
21    dZ1 = np.multiply(dA1, np.int64(A1 > 0))
22    dW1 = 1./m * np.dot(dZ1, X) + (lambda*W1)/m
23    db1 = 1./m * np.sum(dZ1, axis=1, keepdims = True)
24
25
26    gradients = {"dZ4": dZ4, "dW4": dW4, "db4": db4,
27                "dA3": dA3, "dZ3": dZ3, "dW3": dW3, "db3": db3,
28                "dA2": dA2, "dZ2": dZ2, "dW2": dW2, "db2": db2,
29                "dA1": dA1, "dZ1": dZ1, "dW1": dW1, "db1": db1}
30
31    return gradients
```



```
In [148]: 1 def compute_cost_regularisation(A, Y, cache, lambd):
2         #A is predicted
3         #Y is actual
4         m = Y.shape[1]
5         (Z1, A1, W1, b1, Z2, A2, W2, b2, Z3, A3, W3, b3, Z4, A4, W4, b4) = cache
6         cost = compute_cost(A, Y)
7         L2_regularization_cost = (lambd/(2*m))*(np.sum(np.square(W1))+np.sum(np.square(W2))+np.sum(np.square(W3)
8         costTotal = cost + L2_regularization_cost
9         return costTotal
```

```

In [152]: 1 def model_regularise(dataset,numEpochs, lambd):
2         """
3         Model to get the parameters needed to give predicted value of trainset and testset images,
4         final cost and accuracy
5         ---
6         Returns :
7         parameters
8         cost in each iteration
9         accuracy of the model
10
11
12         """
13         k=len(dataset)
14         numBatches=k/batchsize
15         layers_dims = [input_size, hidden_layer_1, hidden_layer_2, hidden_layer_3, output_layer]
16         parameters=initialise_parameter(layers_dims)
17         costs=[]
18         acc=[]
19         for epoch in range(numEpochs):
20             for j in range(int(numBatches)):
21                 # Data loader
22                 loader = torch.utils.data.DataLoader(dataset=dataset,batch_size = batchsize ,shuffle=True)
23                 dataiter = iter(loader)
24                 data = next(dataiter)
25                 X,y = data
26                 X=X.numpy()
27                 y=y.numpy()
28
29                 image_vector_size = 28*28
30                 X = X.reshape(X.shape[0], image_vector_size)      #image is already flattened to x/255
31                 Y=one_hot_encode(y)
32                 y_pred,cache=forward_propagation(X,parameters)
33                 cost=compute_cost_regularisation(y_pred,Y.T,cache, lambd)
34                 gradients=backward_propagation_regularisation(X,Y.T,cache, lambd)
35                 parameters=update_parameters(parameters,gradients,learning_rate)
36                 if j%200 ==0:
37                     print (f'Epoch [{epoch+1}/{numEpochs}], Step [{j+1}/{int(numBatches)}], Loss: {cost.item():
38                             costs.append(cost.item())
39                             acc.append(find_accuracy(Y,y_pred))
40         return parameters, acc, costs

```

```
In [153]: 1 trained_parameters, train_acc, train_costs = model_regularise(train_set, 15, 0.7)
```

```
Epoch [1/15], Step [1/937], Loss: 9.2360
Epoch [1/15], Step [201/937], Loss: 9.1291
Epoch [1/15], Step [401/937], Loss: 8.9900
Epoch [1/15], Step [601/937], Loss: 8.8745
Epoch [1/15], Step [801/937], Loss: 8.8100
Epoch [2/15], Step [1/937], Loss: 8.7460
Epoch [2/15], Step [201/937], Loss: 8.4582
Epoch [2/15], Step [401/937], Loss: 8.2992
Epoch [2/15], Step [601/937], Loss: 8.2440
Epoch [2/15], Step [801/937], Loss: 8.0352
Epoch [3/15], Step [1/937], Loss: 7.9218
Epoch [3/15], Step [201/937], Loss: 7.7341
Epoch [3/15], Step [401/937], Loss: 7.5198
Epoch [3/15], Step [601/937], Loss: 7.7112
Epoch [3/15], Step [801/937], Loss: 7.5202
Epoch [4/15], Step [1/937], Loss: 7.3505
Epoch [4/15], Step [201/937], Loss: 7.4434
Epoch [4/15], Step [401/937], Loss: 6.9799
Epoch [4/15], Step [601/937], Loss: 6.9492
Epoch [4/15], Step [801/937], Loss: 7.0246
Epoch [5/15], Step [1/937], Loss: 6.9824
Epoch [5/15], Step [201/937], Loss: 6.8624
Epoch [5/15], Step [401/937], Loss: 6.7982
Epoch [5/15], Step [601/937], Loss: 6.7315
Epoch [5/15], Step [801/937], Loss: 6.6197
Epoch [6/15], Step [1/937], Loss: 6.6636
Epoch [6/15], Step [201/937], Loss: 6.4167
Epoch [6/15], Step [401/937], Loss: 6.7308
Epoch [6/15], Step [601/937], Loss: 6.5396
Epoch [6/15], Step [801/937], Loss: 6.3591
Epoch [7/15], Step [1/937], Loss: 6.5376
Epoch [7/15], Step [201/937], Loss: 6.3311
Epoch [7/15], Step [401/937], Loss: 6.3302
Epoch [7/15], Step [601/937], Loss: 6.4631
Epoch [7/15], Step [801/937], Loss: 6.2552
Epoch [8/15], Step [1/937], Loss: 6.2491
Epoch [8/15], Step [201/937], Loss: 6.1343
Epoch [8/15], Step [401/937], Loss: 6.4113
Epoch [8/15], Step [601/937], Loss: 6.1886
```

```
Epoch [8/15], Step [801/937], Loss: 6.1422
Epoch [9/15], Step [1/937], Loss: 6.2295
Epoch [9/15], Step [201/937], Loss: 6.1731
Epoch [9/15], Step [401/937], Loss: 5.9459
Epoch [9/15], Step [601/937], Loss: 6.1010
Epoch [9/15], Step [801/937], Loss: 5.9410
Epoch [10/15], Step [1/937], Loss: 6.1796
Epoch [10/15], Step [201/937], Loss: 5.8334
Epoch [10/15], Step [401/937], Loss: 6.2310
Epoch [10/15], Step [601/937], Loss: 6.0370
Epoch [10/15], Step [801/937], Loss: 5.9657
Epoch [11/15], Step [1/937], Loss: 6.0111
Epoch [11/15], Step [201/937], Loss: 5.6325
Epoch [11/15], Step [401/937], Loss: 6.0231
Epoch [11/15], Step [601/937], Loss: 5.6637
Epoch [11/15], Step [801/937], Loss: 5.8209
Epoch [12/15], Step [1/937], Loss: 5.6800
Epoch [12/15], Step [201/937], Loss: 5.9176
Epoch [12/15], Step [401/937], Loss: 5.8904
Epoch [12/15], Step [601/937], Loss: 5.7706
Epoch [12/15], Step [801/937], Loss: 5.7101
Epoch [13/15], Step [1/937], Loss: 5.6675
Epoch [13/15], Step [201/937], Loss: 5.6694
Epoch [13/15], Step [401/937], Loss: 5.9339
Epoch [13/15], Step [601/937], Loss: 5.6345
Epoch [13/15], Step [801/937], Loss: 5.6376
Epoch [14/15], Step [1/937], Loss: 5.7723
Epoch [14/15], Step [201/937], Loss: 5.7646
Epoch [14/15], Step [401/937], Loss: 5.4772
Epoch [14/15], Step [601/937], Loss: 5.4760
Epoch [14/15], Step [801/937], Loss: 5.7334
Epoch [15/15], Step [1/937], Loss: 5.4612
Epoch [15/15], Step [201/937], Loss: 5.4519
Epoch [15/15], Step [401/937], Loss: 5.3695
Epoch [15/15], Step [601/937], Loss: 5.4580
Epoch [15/15], Step [801/937], Loss: 5.5179
```

```
In [154]: 1 print(f'Train Accuracy is : {train_acc[-1]}')
```

```
Train Accuracy is : 0.90625
```

```
In [155]: 1 print ("On the TEST set:")
          2 y_actual_test, y_pred_test = Accuracy(test_set, trained_parameters, 10000)
```

On the TEST set:

accuracy is =0.9022999999999999

/home/mansi/.conda/envs/hbp/lib/python3.7/site-packages/ipykernel_launcher.py:9: RuntimeWarning: overflow encountered in exp

if __name__ == "__main__":

/home/mansi/.conda/envs/hbp/lib/python3.7/site-packages/ipykernel_launcher.py:9: RuntimeWarning: invalid value encountered in true_divide


if __name__ == "__main__":

```

In [156]: 1 #to get the heatmap for the confusion matrix
          2 import seaborn as sn
          3 plt.figure(figsize=(8,8))
          4 sn.heatmap(confusion_matrix(y_actual_test, y_pred_test),annot=True,fmt='d')
          5 plt.xlabel('Predicted Label')
          6 plt.ylabel('Trut Label')
          7 plt.title('Confusion Matrix of Sigmoid Test Data')

```

Out[156]: Text(0.5, 1.0, 'Confusion Matrix of Sigmoid Test Data')



```
In [157]: 1 c = confusion_matrix(y_actual_test, y_pred_test)
          2 print(c)
```

```
[[ 970    0    1    2    1    3    3    0    0    0]
 [    0 1114    4    2    1    1    3    2    8    0]
 [ 110   123  742   10   19    0    8    6   13    1]
 [   62    3  191  727    2    5    1    9    8    2]
 [    2    3   12    1  945    1    6    0    6    6]
 [ 101    2   12  252   25  464    9    4   21    2]
 [ 137    3  309    1  144   34  325    1    4    0]
 [   33   20   60  198   26    6    0  668    6   11]
 [   57   20   86  140   33  193   19   13  412    1]
 [   25    5    6   24  728   14    0  109   20   78]]
```

Observations:

For regularisation output:

- **For $\eta = 0.01$, batch size = 64 and epochs = 15**
 - Train Accuracy = 90.62%
 - Test Accuracy = 90.22%