

PA4_EE21S063

October 30, 2022

0.1 MANSI KAKKAR (EE21S063)

1 PROGRAMMING ASSIGNMENT 4 OF DEEP LEARNING FOR IMAGING (EE5179)

2 Autoencoders

2.1 Libraries Used

```
[64]: import sys
import numpy as np
import os
import matplotlib.pyplot as plt
import torch
import random
from torchvision import datasets
import torchvision.transforms as transforms
from torch.utils.data import Dataset, DataLoader, random_split
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision.utils import make_grid
from sklearn.model_selection import train_test_split
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error as mse
from PIL import Image

torch.manual_seed(2111)

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 10.0) # set default size of plots
```

2.2 Setting up device

```
[2]: #Setting the device to GPU
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)
```

cpu

2.3 Initialisation

```
[3]: batch_size = 500
data_ind = [1000, 6, 375, 789, 150, 771, 4150, 6115, 300, 8141]
epochs = 10
```

2.4 Loading the data

```
[4]: transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.
    ↪1307,), (0.3081,)),])
dataset = datasets.MNIST(root = "data/", train = True, transform = transform,
    ↪download = True)
test_dataset = datasets.MNIST(root = "data/", train = False, transform =
    ↪transform, download = True)

#Splitting the Training dataset
train_dataset, validation_dataset = train_test_split(dataset, test_size=10000,
    ↪train_size = 50000, random_state = None, shuffle = True)

print(f"number of train samples: {len(train_dataset)}")
print(f"number of validation samples: {len(validation_dataset)}")
print(f"number of test samples: {len(test_dataset)}")

train_loader = DataLoader(train_dataset, batch_size = batch_size)
validation_loader = DataLoader(validation_dataset, batch_size = batch_size)
test_loader = DataLoader(test_dataset, batch_size = batch_size)
```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz>
Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to
data/MNIST/raw/train-images-idx3-ubyte.gz

0%| | 0/9912422 [00:00<?, ?it/s]

Extracting data/MNIST/raw/train-images-idx3-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz>

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz> to
data/MNIST/raw/train-labels-idx1-ubyte.gz

0%| | 0/28881 [00:00<?, ?it/s]

Extracting data/MNIST/raw/train-labels-idx1-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz>

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz> to
data/MNIST/raw/t10k-images-idx3-ubyte.gz

0%| | 0/1648877 [00:00<?, ?it/s]

Extracting data/MNIST/raw/t10k-images-idx3-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz>

Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz> to
data/MNIST/raw/t10k-labels-idx1-ubyte.gz

0%| | 0/4542 [00:00<?, ?it/s]

Extracting data/MNIST/raw/t10k-labels-idx1-ubyte.gz to data/MNIST/raw

number of train samples: 50000

number of validation samples: 10000

number of test samples: 10000

```
[5]: print(f"size of train dataloader is :{len(train_loader)}")
      print(f"size of validation dataloader is :{len(validation_loader)}")
      print(f"size of test dataloader is :{len(test_loader)}")
      data = next(iter(train_loader))
      train_img, train_target = data
      print(f"Train image shape:{train_img.shape}")
      print(f"Train Targets shape:{train_target.shape}")

      data = next(iter(validation_loader))
      val_img, val_target = data
      print(f"Validation image shape:{val_img.shape}")
      print(f"Validation Targets shape:{val_target.shape}")

      data = next(iter(test_loader))
      test_img, test_target = data
      print(f"Test image shape:{test_img.shape}")
      print(f"Test Targets shape:{test_target.shape}")
```

size of train dataloader is :100

size of validation dataloader is :20

size of test dataloader is :20

```

Train image shape:torch.Size([500, 1, 28, 28])
Train Targets shape:torch.Size([500])
Validation image shape:torch.Size([500, 1, 28, 28])
Validation Targets shape:torch.Size([500])
Test image shape:torch.Size([500, 1, 28, 28])
Test Targets shape:torch.Size([500])

```

3 1. Comparing PCA and Autoencoders

3.1 PCA

```

[6]: def visualise_data(data, data_ind = data_ind):
      fig, ax = plt.subplots(1, 10, figsize = (20,20))
      for i,ind in enumerate(data_ind):
          image_sample = np.asarray(data[ind],dtype = np.uint8).reshape(28,28)
          image_obj = Image.fromarray(image_sample)
          print('for digit '+str(i))
          ax[i].imshow(image_obj)
      plt.show()

```

```

[7]: pca_train = (dataset.data).reshape(60000,784)
      pca_test = (test_dataset.data).reshape(10000,784)
      data = torch.cat((pca_train, pca_test))
      print(data.shape)
      visualise_data(data)


```

```
torch.Size([70000, 784])
```

```

for digit 0
for digit 1
for digit 2
for digit 3
for digit 4
for digit 5
for digit 6
for digit 7
for digit 8
for digit 9

```




```
[8]: pca_data = data.reshape(70000, 784)
pca_data = np.asarray(pca_data)/255
pca_30 = PCA(n_components = 30)
pca_30.fit(pca_data)
train_pca = pca_30.transform(pca_data)
recover_data = pca_30.inverse_transform(train_pca)
pca_error = mse(pca_data, recover_data)
print(f'reconstruction error is : {pca_error}')

visualise_data(recover_data*255)
```

reconstruction error is : 0.018056392603431226

```
for digit 0
for digit 1
for digit 2
for digit 3
for digit 4
for digit 5
for digit 6
for digit 7
for digit 8
for digit 9
```



```
[9]: # def PCA_implementation(data, num_of_values = 30):
#     mean = torch.mean(data.float(),0)
#     data = data - mean.expand_as(data)
#     U,S,V = torch.svd(torch.mm(data,torch.t(data)))
#     return torch.mm(data,U[:, :num_of_values]),U[:, :num_of_values]
```

```
[10]: # recover_data = PCA_implementation(data, 30)
#Visualisation
# for i,ind in enumerate(data_ind):
#     image_sample = np.asarray(255*recover_data[ind],dtype = np.uint8).
#         ↪ reshape(28,28)
#     #creating image object
#     image_obj = Image.fromarray(image_sample)
#     print('Digit '+str(i))
#     plt.imshow(image_obj)
#     plt.show()
```

3.2 Training Function

```
[47]: def Train(model, optimizer, criterion, epochs = 10, sparse = False, l1_reg = 0.
      ↪001, denoise = False, noise_val = 0.1, model_flag = 0):
    model.train()
    training_loss = []
    validation_loss = []
    # validation_acc = []
    for epoch in range(epochs):
        # val_correct = 0
        for i, (images, labels) in enumerate(train_loader):
            images = images.to(device)
            labels = labels.to(device)
            if(model_flag == 0):
                images = images.reshape(batch_size, 784)

            if(denoise == True):
                noisy = Add_Noise(images, noise_val)
                output, encoded = model(noisy)
            else:
                output, encoded = model(images)

            t_loss = criterion(output, images)

            if(sparse == True):
                t_loss += l1_reg*torch.linalg.norm(encoded,1)

            optimizer.zero_grad()
            t_loss.backward()
            optimizer.step()

        training_loss.append(t_loss)

        for i, (images, labels) in enumerate(validation_loader):
            images = images.to(device)
            labels = labels.to(device)

            if(model_flag==0):
                images = images.reshape(batch_size, 784)

            output, encoded = model(images)
            v_loss = criterion(output, images)
            if(sparse == True):
                v_loss += l1_reg*torch.linalg.norm(encoded,1)
        validation_loss.append(v_loss)
        # validation_acc.append(val_correct/500)
```

```

    print('Epochs: {}/{} ||| with Training Loss = {} ||| Validation Loss = {}'.
    ↪format(epoch+1, epochs, t_loss, v_loss))
    return training_loss, validation_loss

```

3.3 Functions to Visualise Loss and Accuracy

```

[13]: def visualise_loss(trainingloss, valloss):
    with torch.no_grad():
        plt.figure()
        xloss=np.arange(len(trainingloss))
        plt.plot(xloss, trainingloss, color='Blue',label='Training Loss')
        plt.plot(xloss, valloss, color='Orange',label='Validation Loss')
        plt.grid(b=True, which='major', color='#666666', linestyle='-')
        plt.minorticks_on()
        plt.grid(b=True, which='minor', color='#999999', linestyle='--', alpha=0.2)
        plt.legend()
        plt.xlabel('Iterations')
        plt.ylabel(' Loss')
        plt.title(' Loss vs epochs')

def visualise_accuracy(loss_type, loss_list, accuracy):
    with torch.no_grad():
        plt.figure()
        xloss=np.arange(len(loss_list))
        plt.plot(xloss, accuracy)
        plt.grid(b=True, which='major', color='#666666', linestyle='-')
        plt.minorticks_on()
        plt.grid(b=True, which='minor', color='#999999', linestyle='--', alpha=0.2)
        plt.xlabel('Epochs')
        plt.ylabel(loss_type + 'Accuracy')
        plt.title('accuracy vs epochs')

```

3.4 Function to Visualise model

```

[50]: def visualise_model(model, data_ind, data, model_flag = 0):
    with torch.no_grad():
        fig, ax = plt.subplots(1, 10, figsize = (20,20))
        for i,ind in enumerate(data_ind):
            if(model_flag==0):
                output, encoded = model(data[ind].reshape(1,784)).#.cuda())
            else:
                output, encoded = model(data[ind].reshape(1, 28, 28))
            image_sample = np.asarray(output, dtype = np.uint8).reshape(28,28)  #.
            ↪detach().cpu()

```

```

        image_obj = Image.fromarray(image_sample)
        print('for digit '+str(i))
        ax[i].imshow(image_obj)
    plt.show()

```

3.5 Autoencoder

```

[15]: class autoencoder(nn.Module):
        def __init__(self):
            super(autoencoder, self).__init__()
            self.encoder = nn.Sequential(nn.Linear(784, 512), nn.ReLU(), nn.Linear(512, 256), nn.ReLU(), nn.Linear(256, 128), nn.ReLU(), nn.Linear(128, 30), nn.ReLU())
            self.decoder = nn.Sequential(nn.Linear(30, 128), nn.ReLU(), nn.Linear(128, 256), nn.ReLU(), nn.Linear(256, 784), nn.ReLU())
        def forward(self, x):
            encoded = self.encoder(x.float())
            out = self.decoder(encoded)
            return out, encoded

```

Calling Autoencoder function

```

[16]: epochs = 15
        model1 = autoencoder().to(device)
        optimizer1 = optim.Adam(model1.parameters(), lr = 0.003)
        criterion = nn.MSELoss()
        training_loss, val_loss = Train(model1, optimizer1, criterion, epochs)

```


```

Epochs: 1/15 ||| with Training Loss = 0.6887065768241882 ||| Validation Loss = 0.6920874714851379
Epochs: 2/15 ||| with Training Loss = 0.5373498201370239 ||| Validation Loss = 0.5356813073158264
Epochs: 3/15 ||| with Training Loss = 0.483600914478302 ||| Validation Loss = 0.4819205701351166
Epochs: 4/15 ||| with Training Loss = 0.4582567811012268 ||| Validation Loss = 0.45659059286117554
Epochs: 5/15 ||| with Training Loss = 0.4338372051715851 ||| Validation Loss = 0.43202462792396545
Epochs: 6/15 ||| with Training Loss = 0.4201061725616455 ||| Validation Loss = 0.42222604155540466
Epochs: 7/15 ||| with Training Loss = 0.41564664244651794 ||| Validation Loss = 0.4123659133911133
Epochs: 8/15 ||| with Training Loss = 0.405624121427536 ||| Validation Loss = 0.4074176549911499
Epochs: 9/15 ||| with Training Loss = 0.39755842089653015 ||| Validation Loss = 0.39963024854660034

```



Epochs: 10/15 ||| with Training Loss = 0.3911236822605133 ||| Validation Loss = 0.3926449716091156
Epochs: 11/15 ||| with Training Loss = 0.38827431201934814 ||| Validation Loss = 0.3935399055480957
Epochs: 12/15 ||| with Training Loss = 0.38678452372550964 ||| Validation Loss = 0.3860878050327301
Epochs: 13/15 ||| with Training Loss = 0.38392937183380127 ||| Validation Loss = 0.3847467005252838
Epochs: 14/15 ||| with Training Loss = 0.38020917773246765 ||| Validation Loss = 0.3812178373336792
Epochs: 15/15 ||| with Training Loss = 0.3743497431278229 ||| Validation Loss = 0.3766360878944397

```
[17]: visualise_loss(training_loss, val_loss)
```



```
[51]: visualise_model(model1, data_ind, data)
```


```
for digit 0
for digit 1
for digit 2
for digit 3
for digit 4
for digit 5
for digit 6
for digit 7
for digit 8
for digit 9
```



Visualise the encoded output

```
[29]: with torch.no_grad():
      output1, encoded1=model1.forward(data[data_ind[3]].reshape(1,784).float())
      print(encoded1.shape)
      plt.imshow(encoded1.reshape(6, 5))
```

```
torch.Size([1, 30])
```



3.6 Observations:

- From the PCA output we can still interpret the digits but its not neatly displayed(includes a lot of noise in the signal)
- Whereas in the standard Autoencoder output, digits are more clearly visible and interpretable.

- Loss decreases to considerable amount for both training as well as validation

4 2. Experimenting with hidden units of varying sizes



```
[30]: class autoencoder_hidden(nn.Module):
    def __init__(self, hidden_size):
        super(autoencoder_hidden, self).__init__()
        self.encoder = nn.Sequential(nn.Linear(784, hidden_size), nn.ReLU())
        self.decoder = nn.Sequential(nn.Linear(hidden_size, 784), nn.ReLU())

    def forward(self, x):
        encoded = self.encoder(x.float())
        out = self.decoder(encoded)
        return out, encoded
```



```
[63]: epochs = 15
Xnp=np.random.normal(0.5,0.05,(1,28,28))
X=torch.from_numpy(Xnp)
hidden_size = [64, 128, 256]
for i in range(3):
    print("for Hidden Size = {}".format(hidden_size[i]))
    model2 = autoencoder_hidden(hidden_size[i]).to(device)
    optimizer2 = optim.Adam(model2.parameters(), lr = 0.003)
    criterion = nn.MSELoss()
    training_loss, val_loss = Train(model2, optimizer2, criterion, epochs)
    visualise_loss(training_loss, val_loss)
    visualise_model(model2, data_ind, data)
    with torch.no_grad():
        output1=model2(X.reshape(1,784).float())
        plt.figure(1)
        fig,ax = plt.subplots()
        ax=plt.subplot(2,1,1)
        ax.set_xticks([])
        ax.set_yticks([])
        im=ax.imshow(X.detach().numpy()[0],cmap='gray')
        fig,ax = plt.subplots()
        ax=plt.subplot(2,1,2)
        ax.set_xticks([])
        ax.set_yticks([])
        im=ax.imshow(output1[0].reshape(28,28),cmap='gray')
        plt.show()
```

```
for Hidden Size = 64
Epochs: 1/15 ||| with Training Loss = 0.35872408747673035 ||| Validation Loss = 0.3569408357143402
Epochs: 2/15 ||| with Training Loss = 0.3178459703922272 ||| Validation Loss =
```

0.31832829117774963
Epochs: 3/15 ||| with Training Loss = 0.30561742186546326 ||| Validation Loss = 0.3055766224861145
Epochs: 4/15 ||| with Training Loss = 0.30175521969795227 ||| Validation Loss = 0.30178868770599365
Epochs: 5/15 ||| with Training Loss = 0.29709362983703613 ||| Validation Loss = 0.2964833974838257
Epochs: 6/15 ||| with Training Loss = 0.2916001081466675 ||| Validation Loss = 0.29153499007225037
Epochs: 7/15 ||| with Training Loss = 0.2888093888759613 ||| Validation Loss = 0.2889271080493927
Epochs: 8/15 ||| with Training Loss = 0.2881960868835449 ||| Validation Loss = 0.2883595824241638
Epochs: 9/15 ||| with Training Loss = 0.28784826397895813 ||| Validation Loss = 0.2880071997642517
Epochs: 10/15 ||| with Training Loss = 0.28757244348526 ||| Validation Loss = 0.2877322733402252
Epochs: 11/15 ||| with Training Loss = 0.2873050570487976 ||| Validation Loss = 0.2875053882598877
Epochs: 12/15 ||| with Training Loss = 0.28708416223526 ||| Validation Loss = 0.28734394907951355
Epochs: 13/15 ||| with Training Loss = 0.2868974804878235 ||| Validation Loss = 0.2872314751148224
Epochs: 14/15 ||| with Training Loss = 0.28672245144844055 ||| Validation Loss = 0.2871205806732178
Epochs: 15/15 ||| with Training Loss = 0.2859320044517517 ||| Validation Loss = 0.28645116090774536
for digit 0
for digit 1
for digit 2
for digit 3
for digit 4
for digit 5
for digit 6
for digit 7
for digit 8
for digit 9





<Figure size 720x720 with 0 Axes>





for Hidden Size = 128
Epochs: 1/15 ||| with Training Loss = 0.29094651341438293 ||| Validation Loss =

0.2902321219444275
Epochs: 2/15 ||| with Training Loss = 0.2623888850212097 ||| Validation Loss = 0.26192641258239746
Epochs: 3/15 ||| with Training Loss = 0.25273483991622925 ||| Validation Loss = 0.25269418954849243
Epochs: 4/15 ||| with Training Loss = 0.2493942528963089 ||| Validation Loss = 0.24966394901275635
Epochs: 5/15 ||| with Training Loss = 0.24694737792015076 ||| Validation Loss = 0.24729479849338531
Epochs: 6/15 ||| with Training Loss = 0.2447780966758728 ||| Validation Loss = 0.2451203614473343
Epochs: 7/15 ||| with Training Loss = 0.2436942160129547 ||| Validation Loss = 0.2442280799150467
Epochs: 8/15 ||| with Training Loss = 0.24288320541381836 ||| Validation Loss = 0.24324625730514526
Epochs: 9/15 ||| with Training Loss = 0.24240043759346008 ||| Validation Loss = 0.242728590965271
Epochs: 10/15 ||| with Training Loss = 0.24145545065402985 ||| Validation Loss = 0.24192163348197937
Epochs: 11/15 ||| with Training Loss = 0.24027711153030396 ||| Validation Loss = 0.2410450428724289
Epochs: 12/15 ||| with Training Loss = 0.23986749351024628 ||| Validation Loss = 0.24063895642757416
Epochs: 13/15 ||| with Training Loss = 0.23964428901672363 ||| Validation Loss = 0.24045178294181824
Epochs: 14/15 ||| with Training Loss = 0.2390761375427246 ||| Validation Loss = 0.2401057481765747
Epochs: 15/15 ||| with Training Loss = 0.23840881884098053 ||| Validation Loss = 0.23936542868614197
for digit 0
for digit 1
for digit 2
for digit 3
for digit 4
for digit 5
for digit 6
for digit 7
for digit 8
for digit 9





<Figure size 720x720 with 0 Axes>





for Hidden Size = 256
Epochs: 1/15 ||| with Training Loss = 0.2579302191734314 ||| Validation Loss =

0.25769492983818054
Epochs: 2/15 ||| with Training Loss = 0.23645934462547302 ||| Validation Loss = 0.23748131096363068
Epochs: 3/15 ||| with Training Loss = 0.2301545888185501 ||| Validation Loss = 0.23136889934539795
Epochs: 4/15 ||| with Training Loss = 0.22615091502666473 ||| Validation Loss = 0.22747677564620972
Epochs: 5/15 ||| with Training Loss = 0.22367747128009796 ||| Validation Loss = 0.2248115837574005
Epochs: 6/15 ||| with Training Loss = 0.22237077355384827 ||| Validation Loss = 0.22324234247207642
Epochs: 7/15 ||| with Training Loss = 0.22127237915992737 ||| Validation Loss = 0.22221161425113678
Epochs: 8/15 ||| with Training Loss = 0.2205965220928192 ||| Validation Loss = 0.2214614748954773
Epochs: 9/15 ||| with Training Loss = 0.22031590342521667 ||| Validation Loss = 0.22110092639923096
Epochs: 10/15 ||| with Training Loss = 0.22006572782993317 ||| Validation Loss = 0.22084935009479523
Epochs: 11/15 ||| with Training Loss = 0.2199128270149231 ||| Validation Loss = 0.2207784801721573
Epochs: 12/15 ||| with Training Loss = 0.21973708271980286 ||| Validation Loss = 0.2205630987882614
Epochs: 13/15 ||| with Training Loss = 0.2194749265909195 ||| Validation Loss = 0.22037877142429352
Epochs: 14/15 ||| with Training Loss = 0.21941041946411133 ||| Validation Loss = 0.22045297920703888
Epochs: 15/15 ||| with Training Loss = 0.21944372355937958 ||| Validation Loss = 0.22087204456329346
for digit 0
for digit 1
for digit 2
for digit 3
for digit 4
for digit 5
for digit 6
for digit 7
for digit 8
for digit 9



<Figure size 720x720 with 0 Axes>



```
[43]: print("Average Hidden layer activations for 10 images for the hidden layer =  
      ↪256")
```


```

with torch.no_grad():
    sum = 0
    for i in range(10):


        output2, encoded2=model2.forward(data[data_ind[i]].reshape(1,784).float())
        avg = torch.norm(encoded2, p=1)/256.0
        print(f'for the digit = {i} average value is = {avg.detach().numpy()}')
        sum+=avg.detach().numpy()
    plt.imshow(encoded2.reshape(16, 16))
    plt.show()
    print("Average of these values",sum/10.0)

```


Average Hidden layer activations for 10 images for the hidden layer = 256
for the digit = 0 average value is = 44.89842224121094




for the digit = 1 average value is = 41.207969665527344




for the digit = 2 average value is = 60.484375




for the digit = 3 average value is = 46.20608139038086




for the digit = 4 average value is = 48.670780181884766







for the digit = 6 average value is = 45.25163650512695




for the digit = 7 average value is = 54.53059768676758



for the digit = 8 average value is = 41.96870803833008



for the digit = 9 average value is = 36.63874816894531



Average of these values 46.530976104736325

4.1 Observations:

- We can observe that for hidden layer size = 256 we get the best results as compared to the previous ones.
- We observe the best results for size = 256, by comparing both visually and by the training and validation loss obtained

5 3. Sparse Autoencoders

```
[48]: class autoencoder_sparse(nn.Module):
    def __init__(self):
        super(autoencoder_sparse, self).__init__()
        self.encoder = nn.Sequential(nn.Linear(784, 1024), nn.ReLU())
        self.decoder = nn.Sequential(nn.Linear(1024, 784), nn.ReLU())

    def forward(self, x):
        encoded = self.encoder(x.float())
        out = self.decoder(encoded)
        return out, encoded
```


```
[52]: model3 = autoencoder_sparse().to(device)
optimizer3 = optim.Adam(model3.parameters(), lr = 0.0003)
criterion = nn.MSELoss()
training_loss, val_loss = Train(model3, optimizer3, criterion, epochs, sparse =
↳ True, l1_reg = 1e-7)
visualise_loss(training_loss, val_loss)
visualise_model(model3, data_ind, data)
```


```
Epochs: 1/15 ||| with Training Loss = 0.2669614553451538 ||| Validation Loss =
0.2661038041114807
Epochs: 2/15 ||| with Training Loss = 0.2304380089044571 ||| Validation Loss =
0.23078307509422302
Epochs: 3/15 ||| with Training Loss = 0.2187575101852417 ||| Validation Loss =
0.21954764425754547
Epochs: 4/15 ||| with Training Loss = 0.21263164281845093 ||| Validation Loss =
0.21345174312591553
Epochs: 5/15 ||| with Training Loss = 0.20871004462242126 ||| Validation Loss =
0.2095082700252533
Epochs: 6/15 ||| with Training Loss = 0.20602142810821533 ||| Validation Loss =
0.20674504339694977
Epochs: 7/15 ||| with Training Loss = 0.20398883521556854 ||| Validation Loss =
0.20471002161502838
Epochs: 8/15 ||| with Training Loss = 0.20240412652492523 ||| Validation Loss =
0.20313666760921478
Epochs: 9/15 ||| with Training Loss = 0.20117852091789246 ||| Validation Loss =
0.20189400017261505
Epochs: 10/15 ||| with Training Loss = 0.20018678903579712 ||| Validation Loss =
0.2009064108133316
Epochs: 11/15 ||| with Training Loss = 0.19939512014389038 ||| Validation Loss =
0.20010945200920105
Epochs: 12/15 ||| with Training Loss = 0.1987469643354416 ||| Validation Loss =
0.19946467876434326
Epochs: 13/15 ||| with Training Loss = 0.19822199642658234 ||| Validation Loss =
0.19892361760139465
```


Epochs: 14/15 ||| with Training Loss = 0.19779588282108307 ||| Validation Loss = 0.19846349954605103

Epochs: 15/15 ||| with Training Loss = 0.19742344319820404 ||| Validation Loss = 0.19808447360992432

```
for digit 0
for digit 1
for digit 2
for digit 3
for digit 4
for digit 5
for digit 6
for digit 7
for digit 8
for digit 9
```







```
[54]: print("Average Hidden layer activations for 10 images for the hidden layer = 256")
with torch.no_grad():
    sum = 0
    for i in range(10):

        output3, encoded3=model3.forward(data[data_ind[i]].reshape(1,784).float())
        avg = torch.norm(encoded3, p=1)/1024.0
        print(f'for the digit = {i} average value is = {avg.detach().numpy()}')
        sum+=avg.detach().numpy()
        plt.imshow(encoded3.reshape(32, 32))
        plt.show()
    print("Average of these values",sum/10.0)
```


Average Hidden layer activations for 10 images for the hidden layer = 256
for the digit = 0 average value is = 76.5195083618164




for the digit = 1 average value is = 65.42584991455078




for the digit = 2 average value is = 93.7727279663086




for the digit = 3 average value is = 85.11482238769531




for the digit = 4 average value is = 86.35318756103516




for the digit = 5 average value is = 81.92125701904297




for the digit = 6 average value is = 76.35543823242188




for the digit = 7 average value is = 111.32823181152344



for the digit = 8 average value is = 78.93204498291016



for the digit = 9 average value is = 65.29202270507812




Average of these values 82.10150909423828

```
[55]: ix=1
fig,ax=plt.subplots()
for i in range(len(model3.state_dict()['encoder.0.weight'])):
    ax=plt.subplot(35,35,ix)
    ax.set_xticks([])
    ax.set_yticks([])
    im=ax.imshow(model3.state_dict()['encoder.0.weight'][i].
↪reshape(28,28),cmap='gray')
    ix+=1


fig.subplots_adjust(right=0.8)
```

```
cbar_ax = fig.add_axes([0.85, 0.15, 0.05, 0.7])
fig.colorbar(im, cax=cbar_ax)
plt.show()
```



```
[57]: plt.imshow(model3.state_dict()['encoder.0.weight'], cmap='gray')
```


```
[57]: <matplotlib.image.AxesImage at 0x7fef0ac53d90>
```




```
[71]: with torch.no_grad():  
      a=random.sample(range(0,784),int(0.9*784))  
      X=data[data_ind[3]].clone()  
      X=X.reshape(1,784)  
      X[0][a]=0
```

```
output1=model3(data[data_ind[3]].reshape(1,784))
plt.imshow(output1[0].reshape(28,28),cmap='gray')
plt.show()

output2=model3(X.reshape(1,784))
plt.imshow(output2[0].reshape(28,28),cmap='gray')
plt.show()
```





6 4. Denoising Encoder

```
[72]: def Add_Noise(image, noise_val = 0.3):  
      noise = torch.randn(image.size())*noise_val  
      noisy_image = image + noise  
      return noisy_image
```


6.1 First using autoencoder in question 2

```
[73]: noise_val = [0.3, 0.5, 0.8, 0.9]
      for i in range(4):
          print(f'\n for noise value = {noise_val[i]}\n')
          model4 = autoencoder_hidden(256).to(device)
          optimizer4 = optim.Adam(model4.parameters(), lr = 0.003)
          criterion = nn.MSELoss()
          training_loss, val_loss = Train(model4, optimizer4, criterion, epochs,
          ↪denoise = True, noise_val = noise_val[i])

          visualise_loss(training_loss, val_loss)
          visualise_model(model4, data_ind, data)
```



for noise value = 0.3

```
Epochs: 1/15 ||| with Training Loss = 0.25964343547821045 ||| Validation Loss =
0.2541666626930237
Epochs: 2/15 ||| with Training Loss = 0.2404690831899643 ||| Validation Loss =
0.23458996415138245
Epochs: 3/15 ||| with Training Loss = 0.2332722246646881 ||| Validation Loss =
0.2272350937128067
Epochs: 4/15 ||| with Training Loss = 0.2301328331232071 ||| Validation Loss =
0.22372418642044067
Epochs: 5/15 ||| with Training Loss = 0.2286226749420166 ||| Validation Loss =
0.22252203524112701
Epochs: 6/15 ||| with Training Loss = 0.22831477224826813 ||| Validation Loss =
0.22136977314949036
Epochs: 7/15 ||| with Training Loss = 0.22721241414546967 ||| Validation Loss =
0.2199418842792511
Epochs: 8/15 ||| with Training Loss = 0.2257051020860672 ||| Validation Loss =
0.2188408076763153
Epochs: 9/15 ||| with Training Loss = 0.22552256286144257 ||| Validation Loss =
0.2185373455286026
Epochs: 10/15 ||| with Training Loss = 0.22561019659042358 ||| Validation Loss =
0.21870316565036774
Epochs: 11/15 ||| with Training Loss = 0.22540491819381714 ||| Validation Loss =
0.21781963109970093
Epochs: 12/15 ||| with Training Loss = 0.22466717660427094 ||| Validation Loss =
0.21759797632694244
Epochs: 13/15 ||| with Training Loss = 0.2250085324048996 ||| Validation Loss =
0.21756429970264435
Epochs: 14/15 ||| with Training Loss = 0.22451135516166687 ||| Validation Loss =
0.21700796484947205
Epochs: 15/15 ||| with Training Loss = 0.22496536374092102 ||| Validation Loss =
0.21740809082984924
```

```

for digit 0
for digit 1
for digit 2
for digit 3
for digit 4
for digit 5
for digit 6
for digit 7
for digit 8
for digit 9

```





```

for noise value = 0.5

Epochs: 1/15 ||| with Training Loss = 0.2878321409225464 ||| Validation Loss =
0.27613434195518494
Epochs: 2/15 ||| with Training Loss = 0.26901760697364807 ||| Validation Loss =
0.25691479444503784
Epochs: 3/15 ||| with Training Loss = 0.26329925656318665 ||| Validation Loss =
0.25007352232933044
Epochs: 4/15 ||| with Training Loss = 0.2602956295013428 ||| Validation Loss =
0.24704982340335846
Epochs: 5/15 ||| with Training Loss = 0.25883233547210693 ||| Validation Loss =
0.244818314909935
Epochs: 6/15 ||| with Training Loss = 0.25793299078941345 ||| Validation Loss =
0.24308356642723083
Epochs: 7/15 ||| with Training Loss = 0.25811874866485596 ||| Validation Loss =
0.24273598194122314
Epochs: 8/15 ||| with Training Loss = 0.25736838579177856 ||| Validation Loss =
0.24233466386795044
Epochs: 9/15 ||| with Training Loss = 0.2568536698818207 ||| Validation Loss =
0.24215535819530487
Epochs: 10/15 ||| with Training Loss = 0.2565525472164154 ||| Validation Loss =
0.24168632924556732
Epochs: 11/15 ||| with Training Loss = 0.25758954882621765 ||| Validation Loss =
0.24195854365825653
Epochs: 12/15 ||| with Training Loss = 0.25664934515953064 ||| Validation Loss =
0.24104103446006775
Epochs: 13/15 ||| with Training Loss = 0.2558158338069916 ||| Validation Loss =
0.2412956953048706
Epochs: 14/15 ||| with Training Loss = 0.2565968334674835 ||| Validation Loss =
0.241303950548172
Epochs: 15/15 ||| with Training Loss = 0.25577935576438904 ||| Validation Loss =
0.24073509871959686
for digit 0
for digit 1
for digit 2
for digit 3
for digit 4
for digit 5
for digit 6
for digit 7
for digit 8
for digit 9

```





for noise value = 0.8

Epochs: 1/15 ||| with Training Loss = 0.29636359214782715 ||| Validation Loss = 0.26738864183425903

Epochs: 2/15 ||| with Training Loss = 0.2816535234451294 ||| Validation Loss = 0.2507992386817932

Epochs: 3/15 ||| with Training Loss = 0.27668309211730957 ||| Validation Loss = 0.24541570246219635
Epochs: 4/15 ||| with Training Loss = 0.275882750749588 ||| Validation Loss = 0.24309852719306946
Epochs: 5/15 ||| with Training Loss = 0.27408668398857117 ||| Validation Loss = 0.24099692702293396
Epochs: 6/15 ||| with Training Loss = 0.2728397846221924 ||| Validation Loss = 0.2408573031425476
Epochs: 7/15 ||| with Training Loss = 0.2736726701259613 ||| Validation Loss = 0.24050545692443848
Epochs: 8/15 ||| with Training Loss = 0.2728123962879181 ||| Validation Loss = 0.23954950273036957
Epochs: 9/15 ||| with Training Loss = 0.2721427381038666 ||| Validation Loss = 0.23900291323661804
Epochs: 10/15 ||| with Training Loss = 0.27132874727249146 ||| Validation Loss = 0.23917683959007263
Epochs: 11/15 ||| with Training Loss = 0.2719590961933136 ||| Validation Loss = 0.2390621155500412
Epochs: 12/15 ||| with Training Loss = 0.2716068625450134 ||| Validation Loss = 0.23901471495628357
Epochs: 13/15 ||| with Training Loss = 0.2719190716743469 ||| Validation Loss = 0.23834188282489777
Epochs: 14/15 ||| with Training Loss = 0.27185115218162537 ||| Validation Loss = 0.2388356328010559
Epochs: 15/15 ||| with Training Loss = 0.2710857689380646 ||| Validation Loss = 0.2379731684923172
for digit 0
for digit 1
for digit 2
for digit 3
for digit 4
for digit 5
for digit 6
for digit 7
for digit 8
for digit 9





for noise value = 0.9

Epochs: 1/15 ||| with Training Loss = 0.3061346113681793 ||| Validation Loss = 0.2713543474674225

Epochs: 2/15 ||| with Training Loss = 0.28712302446365356 ||| Validation Loss = 0.2509884238243103

Epochs: 3/15 ||| with Training Loss = 0.2818964719772339 ||| Validation Loss = 0.24540047347545624
Epochs: 4/15 ||| with Training Loss = 0.2795545160770416 ||| Validation Loss = 0.24252009391784668
Epochs: 5/15 ||| with Training Loss = 0.2771940231323242 ||| Validation Loss = 0.24065808951854706
Epochs: 6/15 ||| with Training Loss = 0.27806445956230164 ||| Validation Loss = 0.24072566628456116
Epochs: 7/15 ||| with Training Loss = 0.2770940065383911 ||| Validation Loss = 0.24046345055103302
Epochs: 8/15 ||| with Training Loss = 0.27676868438720703 ||| Validation Loss = 0.2393503040075302
Epochs: 9/15 ||| with Training Loss = 0.2761901319026947 ||| Validation Loss = 0.23872050642967224
Epochs: 10/15 ||| with Training Loss = 0.27673855423927307 ||| Validation Loss = 0.23862308263778687
Epochs: 11/15 ||| with Training Loss = 0.2765922546386719 ||| Validation Loss = 0.23879817128181458
Epochs: 12/15 ||| with Training Loss = 0.2766468822956085 ||| Validation Loss = 0.2387128472328186
Epochs: 13/15 ||| with Training Loss = 0.2759872376918793 ||| Validation Loss = 0.2384033203125
Epochs: 14/15 ||| with Training Loss = 0.2755905091762543 ||| Validation Loss = 0.2386041283607483
Epochs: 15/15 ||| with Training Loss = 0.27587398886680603 ||| Validation Loss = 0.2379777580499649
for digit 0
for digit 1
for digit 2
for digit 3
for digit 4
for digit 5
for digit 6
for digit 7
for digit 8
for digit 9



6.2 Observations:

- we can observe that the training loss is minimum for the least noise value for the standard autoencoder in question 2


```
[74]: class autoencoder_denoising(nn.Module):
    def __init__(self):
        super(autoencoder_denoising, self).__init__()
        self.encoder = nn.Sequential(nn.Linear(784,64), nn.ReLU(), nn.Linear(64,8),
        ↪nn.ReLU())
        self.decoder =nn.Sequential(nn.Linear(8,64), nn.ReLU(), nn.Linear(64,784),
        ↪nn.ReLU())

    def forward(self,x):
        encoded = self.encoder(x.float())
        out = self.decoder(encoded)
        return out, encoded
```



```
[75]: noise_val = [0.3, 0.5, 0.8, 0.9]
for i in range(4):
    print(f'\n for noise value = {noise_val[i]}\n')
    model5 = autoencoder_denoising().to(device)
    optimizer5 = optim.Adam(model5.parameters(), lr = 0.003)
    criterion = nn.MSELoss()
    training_loss, val_loss = Train(model5, optimizer5, criterion, epochs,
    ↪denoise = True, noise_val = noise_val[i])

    visualise_loss(training_loss, val_loss)
    visualise_model(model5, data_ind, data)
```

for noise value = 0.3

```
Epochs: 1/15 ||| with Training Loss = 0.6848006248474121 ||| Validation Loss =
0.6831685304641724
Epochs: 2/15 ||| with Training Loss = 0.6406914591789246 ||| Validation Loss =
0.640123724937439
Epochs: 3/15 ||| with Training Loss = 0.6226150989532471 ||| Validation Loss =
0.620271623134613
Epochs: 4/15 ||| with Training Loss = 0.6120909452438354 ||| Validation Loss =
0.6102980375289917
Epochs: 5/15 ||| with Training Loss = 0.6017044186592102 ||| Validation Loss =
0.601205587387085
Epochs: 6/15 ||| with Training Loss = 0.596468985080719 ||| Validation Loss =
0.5963802337646484
Epochs: 7/15 ||| with Training Loss = 0.5918829441070557 ||| Validation Loss =
0.5904800891876221
Epochs: 8/15 ||| with Training Loss = 0.5683847069740295 ||| Validation Loss =
0.5687422156333923
Epochs: 9/15 ||| with Training Loss = 0.5593767166137695 ||| Validation Loss =
0.5602560639381409
Epochs: 10/15 ||| with Training Loss = 0.5549748539924622 ||| Validation Loss =
```

```
0.5555382370948792
Epochs: 11/15 ||| with Training Loss = 0.5505425930023193 ||| Validation Loss =
0.5525535941123962
Epochs: 12/15 ||| with Training Loss = 0.5483794212341309 ||| Validation Loss =
0.550687313079834
Epochs: 13/15 ||| with Training Loss = 0.5464069247245789 ||| Validation Loss =
0.5487973093986511
Epochs: 14/15 ||| with Training Loss = 0.5394445657730103 ||| Validation Loss =
0.5409221053123474
Epochs: 15/15 ||| with Training Loss = 0.5278555750846863 ||| Validation Loss =
0.5303293466567993
for digit 0
for digit 1
for digit 2
for digit 3
for digit 4
for digit 5
for digit 6
for digit 7
for digit 8
for digit 9
```





for noise value = 0.5

Epochs: 1/15 ||| with Training Loss = 0.6872881054878235 ||| Validation Loss = 0.6823856234550476

Epochs: 2/15 ||| with Training Loss = 0.6112332344055176 ||| Validation Loss = 0.6048036217689514

Epochs: 3/15 ||| with Training Loss = 0.5865513682365417 ||| Validation Loss = 0.5807957053184509
Epochs: 4/15 ||| with Training Loss = 0.5761266350746155 ||| Validation Loss = 0.5716156959533691
Epochs: 5/15 ||| with Training Loss = 0.5708513855934143 ||| Validation Loss = 0.5662994384765625
Epochs: 6/15 ||| with Training Loss = 0.5661503672599792 ||| Validation Loss = 0.5615786910057068
Epochs: 7/15 ||| with Training Loss = 0.5633235573768616 ||| Validation Loss = 0.5578161478042603
Epochs: 8/15 ||| with Training Loss = 0.5487060546875 ||| Validation Loss = 0.5449833273887634
Epochs: 9/15 ||| with Training Loss = 0.5392489433288574 ||| Validation Loss = 0.5347309708595276
Epochs: 10/15 ||| with Training Loss = 0.5357452034950256 ||| Validation Loss = 0.5301650166511536
Epochs: 11/15 ||| with Training Loss = 0.5321444272994995 ||| Validation Loss = 0.5256530046463013
Epochs: 12/15 ||| with Training Loss = 0.5295639634132385 ||| Validation Loss = 0.522879958152771
Epochs: 13/15 ||| with Training Loss = 0.5279529690742493 ||| Validation Loss = 0.5204227566719055
Epochs: 14/15 ||| with Training Loss = 0.5261601805686951 ||| Validation Loss = 0.5184678435325623
Epochs: 15/15 ||| with Training Loss = 0.5252169370651245 ||| Validation Loss = 0.5170124769210815
for digit 0
for digit 1
for digit 2
for digit 3
for digit 4
for digit 5
for digit 6
for digit 7
for digit 8
for digit 9





for noise value = 0.8

Epochs: 1/15 ||| with Training Loss = 0.5721724033355713 ||| Validation Loss = 0.5582663416862488

Epochs: 2/15 ||| with Training Loss = 0.5323060154914856 ||| Validation Loss = 0.5217342972755432

Epochs: 3/15 ||| with Training Loss = 0.5184629559516907 ||| Validation Loss =

```
0.5045705437660217
Epochs: 4/15 ||| with Training Loss = 0.5099622011184692 ||| Validation Loss =
0.4970572292804718
Epochs: 5/15 ||| with Training Loss = 0.5041430592536926 ||| Validation Loss =
0.49176591634750366
Epochs: 6/15 ||| with Training Loss = 0.4995822310447693 ||| Validation Loss =
0.4867702126502991
Epochs: 7/15 ||| with Training Loss = 0.49537867307662964 ||| Validation Loss =
0.4835708737373352
Epochs: 8/15 ||| with Training Loss = 0.49420687556266785 ||| Validation Loss =
0.4820214807987213
Epochs: 9/15 ||| with Training Loss = 0.4906397759914398 ||| Validation Loss =
0.47901052236557007
Epochs: 10/15 ||| with Training Loss = 0.4886122941970825 ||| Validation Loss =
0.47735312581062317
Epochs: 11/15 ||| with Training Loss = 0.48786231875419617 ||| Validation Loss =
0.475845068693161
Epochs: 12/15 ||| with Training Loss = 0.4866437315940857 ||| Validation Loss =
0.4749729633331299
Epochs: 13/15 ||| with Training Loss = 0.48546794056892395 ||| Validation Loss =
0.47420862317085266
Epochs: 14/15 ||| with Training Loss = 0.4855162501335144 ||| Validation Loss =
0.4730307459831238
Epochs: 15/15 ||| with Training Loss = 0.4832351803779602 ||| Validation Loss =
0.4730837941169739
for digit 0
for digit 1
for digit 2
for digit 3
for digit 4
for digit 5
for digit 6
for digit 7
for digit 8
for digit 9
```





for noise value = 0.9

Epochs: 1/15 ||| with Training Loss = 0.5881263017654419 ||| Validation Loss = 0.5775903463363647


Epochs: 2/15 ||| with Training Loss = 0.5478460192680359 ||| Validation Loss = 0.5411717295646667

Epochs: 3/15 ||| with Training Loss = 0.5340295433998108 ||| Validation Loss = 0.5279478430747986
Epochs: 4/15 ||| with Training Loss = 0.5278148651123047 ||| Validation Loss = 0.5203167200088501
Epochs: 5/15 ||| with Training Loss = 0.5223813652992249 ||| Validation Loss = 0.5154925584793091
Epochs: 6/15 ||| with Training Loss = 0.515595555305481 ||| Validation Loss = 0.508739173412323
Epochs: 7/15 ||| with Training Loss = 0.5119762420654297 ||| Validation Loss = 0.5050249695777893
Epochs: 8/15 ||| with Training Loss = 0.5084444880485535 ||| Validation Loss = 0.5034490823745728
Epochs: 9/15 ||| with Training Loss = 0.5062755346298218 ||| Validation Loss = 0.500933825969696
Epochs: 10/15 ||| with Training Loss = 0.5037600994110107 ||| Validation Loss = 0.4989137053489685
Epochs: 11/15 ||| with Training Loss = 0.4992004930973053 ||| Validation Loss = 0.4941493272781372
Epochs: 12/15 ||| with Training Loss = 0.49911805987358093 ||| Validation Loss = 0.4928869605064392
Epochs: 13/15 ||| with Training Loss = 0.49759629368782043 ||| Validation Loss = 0.49133554100990295
Epochs: 14/15 ||| with Training Loss = 0.4970111548900604 ||| Validation Loss = 0.49001121520996094
Epochs: 15/15 ||| with Training Loss = 0.4951775372028351 ||| Validation Loss = 0.48928576707839966
for digit 0
for digit 1
for digit 2
for digit 3
for digit 4
for digit 5
for digit 6
for digit 7
for digit 8
for digit 9



```
[86]: with torch.no_grad():
      output5, encoded5 = model5.forward(data[data_ind[3]].reshape(1,784).float())
      print(encoded5.shape)
      plt.imshow(encoded5.reshape(4, 2))
```

```
torch.Size([1, 8])
```



7 Convolutional Autoencoders

Unpooling Convolutional Autoencoder

Encoder Module: * 28X 28 X 1 to 14 X 14 X 8 * 14 X 14 X 8 to 7 X 7 X 16 * 7 X 7 X 16 to 3 X 3 X 16

Decoder Module: * 7 X 7 X 16 to 7 X 7 X 16 * 7 X 7 X 16 to 14 X 14 X 8 * 14 X 14 X 8 to 28 X 28 X 1

```
[76]: class autoencoder_convolutional_unpool(nn.Module):
    def __init__(self):
        super(autoencoder_convolutional_unpool, self).__init__()
        #initializing the encoder module
        self.encoder_conv1 = nn.Sequential(
            nn.Conv2d(1,8, kernel_size = 3, stride = 1,padding= 1),nn.ReLU(),nn.
            ↪MaxPool2d(kernel_size = (2,2),return_indices = True))
        self.encoder_conv2 = nn.Sequential(
            nn.Conv2d(8,16, kernel_size = 3, stride = 1,padding= 1),nn.ReLU(),nn.
            ↪MaxPool2d(kernel_size = (2,2),return_indices = True))
        self.encoder_conv3 = nn.Sequential(
            nn.Conv2d(16,16, kernel_size = 3, stride = 1,padding= 1),nn.ReLU(),nn.
            ↪MaxPool2d(kernel_size = (2,2),return_indices = True))

        #initializing the decoder module
        self.decoder_conv1 = nn.Sequential(
            nn.Identity())
        self.decoder_conv2 = nn.Sequential(
            nn.Conv2d(16,8, kernel_size = 3, stride = 1,padding= 1),nn.ReLU())
        self.decoder_conv3 = nn.Sequential(
            nn.Conv2d(8,1, kernel_size = 3, stride = 1,padding= 1),nn.ReLU())

        #defining the unpooling operation
        self.unpool = nn.MaxUnpool2d(kernel_size = (2,2))

    def forward(self,x): #defines the forward pass and also the structure of the
    ↪network thus helping backprop
        encoded_input,indices1 = self.encoder_conv1(x.float())
        encoded_input,indices2 = self.encoder_conv2(encoded_input)
        encoded_input,indices3 = self.encoder_conv3(encoded_input)

        reconstructed_input = self.
        ↪unpool(encoded_input,indices3,output_size=torch.Size([batch_size, 16, 7, 7]))
        reconstructed_input = self.decoder_conv1(reconstructed_input)
        reconstructed_input = self.unpool(reconstructed_input,indices2)
        reconstructed_input = self.decoder_conv2(reconstructed_input)
```

```

reconstructed_input      = self.unpool(reconstructed_input, indices1)
reconstructed_input      = self.decoder_conv3(reconstructed_input)

return reconstructed_input, encoded_input

```

```


[77]: epochs = 10
model6 = autoencoder_convolutional_unpool().to(device)
optimizer6 = optim.Adam(model6.parameters(), lr = 0.003)
criterion = nn.MSELoss()
training_loss, val_loss = Train(model6, optimizer6, criterion, epochs, ↵
    ↪model_flag=1)
visualise_loss(training_loss, val_loss)

```

```


Epochs: 1/10 ||| with Training Loss = 0.30504223704338074 ||| Validation Loss =
0.3003740906715393
Epochs: 2/10 ||| with Training Loss = 0.26394376158714294 ||| Validation Loss =
0.2607918977737427
Epochs: 3/10 ||| with Training Loss = 0.25145184993743896 ||| Validation Loss =
0.24928033351898193
Epochs: 4/10 ||| with Training Loss = 0.23987926542758942 ||| Validation Loss =
0.23818303644657135
Epochs: 5/10 ||| with Training Loss = 0.22902315855026245 ||| Validation Loss =
0.22745725512504578
Epochs: 6/10 ||| with Training Loss = 0.21915584802627563 ||| Validation Loss =
0.21722395718097687
Epochs: 7/10 ||| with Training Loss = 0.21203015744686127 ||| Validation Loss =
0.2103302925825119
Epochs: 8/10 ||| with Training Loss = 0.2068602591753006 ||| Validation Loss =
0.20536252856254578
Epochs: 9/10 ||| with Training Loss = 0.20335638523101807 ||| Validation Loss =
0.2016081064939499
Epochs: 10/10 ||| with Training Loss = 0.2005344033241272 ||| Validation Loss =
0.19886523485183716

```



```
[78]: visualise_model(model6, data_ind, data, model_flag = 1)
```

```
for digit 0  
for digit 1  
for digit 2  
for digit 3  
for digit 4  
for digit 5  
for digit 6  
for digit 7  
for digit 8  
for digit 9
```



Unpooling + Deconvolution

```
[79]: class autoencoder_deconv_unpool(nn.Module):
    def __init__(self):
        super(autoencoder_deconv_unpool,self).__init__()

        #initializing the encoder module
        self.encoder_conv1 = nn.Sequential(
            nn.Conv2d(1,8, kernel_size = 3, stride = 1,padding= 1),nn.ReLU(),nn.
↪MaxPool2d(kernel_size = (2,2),return_indices = True))
        self.encoder_conv2 = nn.Sequential(
            nn.Conv2d(8,16, kernel_size = 3, stride = 1,padding= 1),nn.ReLU(),nn.
↪MaxPool2d(kernel_size = (2,2),return_indices = True))
        self.encoder_conv3 = nn.Sequential(
            nn.Conv2d(16,16, kernel_size = 3, stride = 1,padding= 1),nn.ReLU(),nn.
↪MaxPool2d(kernel_size = (2,2),return_indices = True))

        #initializing the decoder module
        self.decoder_conv1 = nn.Sequential(
            nn.ConvTranspose2d(16,16, kernel_size = 3, stride = 1, padding = 1),nn.
↪ReLU())
        self.decoder_conv2 = nn.Sequential(
            nn.ConvTranspose2d(16,8, kernel_size = 3, stride = 1, padding = 1),nn.
↪ReLU())
        self.decoder_conv3 = nn.Sequential(
            nn.ConvTranspose2d(8,1, kernel_size = 3, stride = 1, padding = 1),nn.
↪ReLU())

        #defining the unpooling operation
        self.unpool = nn.MaxUnpool2d(kernel_size = (2,2))

    def forward(self,x):

        encoded_input,indices1 = self.encoder_conv1(x.float())
        encoded_input,indices2 = self.encoder_conv2(encoded_input)
        encoded_input,indices3 = self.encoder_conv3(encoded_input)

        reconstructed_input = self.
↪unpool(encoded_input,indices3,output_size=torch.Size([batch_size, 16, 7, 7]))
```

```

reconstructed_input      = self.decoder_conv1(reconstructed_input)
reconstructed_input      = self.unpool(reconstructed_input, indices2)
reconstructed_input      = self.decoder_conv2(reconstructed_input)
reconstructed_input      = self.unpool(reconstructed_input, indices1)
reconstructed_input      = self.decoder_conv3(reconstructed_input)

return reconstructed_input, encoded_input

```

```


[80]: epochs = 10
model8 = autoencoder_deconv_unpool().to(device)
optimizer8 = optim.Adam(model8.parameters(), lr = 0.003)
criterion = nn.MSELoss()
training_loss, val_loss = Train(model8, optimizer8, criterion, epochs,
    ↪model_flag=1)
visualise_loss(training_loss, val_loss)

```

```


Epochs: 1/10 ||| with Training Loss = 0.25539734959602356 ||| Validation Loss =
0.2514868378639221
Epochs: 2/10 ||| with Training Loss = 0.2148585319519043 ||| Validation Loss =
0.2131403684616089
Epochs: 3/10 ||| with Training Loss = 0.2032325565814972 ||| Validation Loss =
0.20196329057216644
Epochs: 4/10 ||| with Training Loss = 0.1985708326101303 ||| Validation Loss =
0.19741162657737732
Epochs: 5/10 ||| with Training Loss = 0.1945735067129135 ||| Validation Loss =
0.19352756440639496
Epochs: 6/10 ||| with Training Loss = 0.1902538686990738 ||| Validation Loss =
0.18954743444919586
Epochs: 7/10 ||| with Training Loss = 0.1876154989004135 ||| Validation Loss =
0.18679484724998474
Epochs: 8/10 ||| with Training Loss = 0.18576161563396454 ||| Validation Loss =
0.18499350547790527
Epochs: 9/10 ||| with Training Loss = 0.18407995998859406 ||| Validation Loss =
0.18318553268909454
Epochs: 10/10 ||| with Training Loss = 0.18284496665000916 ||| Validation Loss =
0.18181955814361572

```



```
[81]: visualise_model(model8, data_ind, data, model_flag = 1)
```

```
for digit 0
for digit 1
for digit 2
for digit 3
for digit 4
for digit 5
for digit 6
for digit 7
for digit 8
for digit 9
```

Deconvolution

```
[82]: class autoencoder_conv_deconv(nn.Module):
    def __init__(self):
        super(autoencoder_conv_deconv,self).__init__()

        #initializing the encoder module
        self.encoder_conv1 = nn.Sequential(
            nn.Conv2d(1,8, kernel_size = 3, stride = 1,padding= 1),nn.ReLU(),nn.
↪MaxPool2d(kernel_size = (2,2)))
        self.encoder_conv2 = nn.Sequential(
            nn.Conv2d(8,16, kernel_size = 3, stride = 1,padding= 1),nn.ReLU(),nn.
↪MaxPool2d(kernel_size = (2,2)))
        self.encoder_conv3 = nn.Sequential(
            nn.Conv2d(16,16, kernel_size = 3, stride = 1,padding= 1),nn.ReLU(),nn.
↪MaxPool2d(kernel_size = (2,2)))

        #initializing the decoder module
        self.decoder_conv1 = nn.Sequential(
            nn.ConvTranspose2d(16,16, kernel_size = 3, stride = 2),nn.ReLU())
        self.decoder_conv2 = nn.Sequential(
            nn.ConvTranspose2d(16,8, kernel_size = 4, stride = 2, padding = 1),nn.
↪ReLU())
        self.decoder_conv3 = nn.Sequential(
            nn.ConvTranspose2d(8,1, kernel_size = 4, stride = 2, padding = 1),nn.
↪ReLU())

    def forward(self,x):


        encoded_input  = self.encoder_conv1(x.float())
        encoded_input  = self.encoder_conv2(encoded_input)
        encoded_input  = self.encoder_conv3(encoded_input)

        reconstructed_input      = self.decoder_conv1(encoded_input)
        reconstructed_input      = self.decoder_conv2(reconstructed_input)
        reconstructed_input      = self.decoder_conv3(reconstructed_input)

        return reconstructed_input,encoded_input
```


```
[83]: epochs = 10
      model7 = autoencoder_conv_deconv().to(device)
      optimizer7 = optim.Adam(model7.parameters(), lr = 0.003)
      criterion = nn.MSELoss()
      training_loss, val_loss = Train(model7, optimizer7, criterion, epochs,
      ↪model_flag=1)
      visualise_loss(training_loss, val_loss)
```

```
Epochs: 1/10 ||| with Training Loss = 0.42669978737831116 ||| Validation Loss =
0.41601505875587463
Epochs: 2/10 ||| with Training Loss = 0.34065744280815125 ||| Validation Loss =
0.3357812762260437
Epochs: 3/10 ||| with Training Loss = 0.30736157298088074 ||| Validation Loss =
0.3067431151866913
Epochs: 4/10 ||| with Training Loss = 0.28730306029319763 ||| Validation Loss =
0.28645655512809753
Epochs: 5/10 ||| with Training Loss = 0.27460142970085144 ||| Validation Loss =
0.2739414572715759
Epochs: 6/10 ||| with Training Loss = 0.2669924795627594 ||| Validation Loss =
0.266570508480072
Epochs: 7/10 ||| with Training Loss = 0.260661780834198 ||| Validation Loss =
0.25891631841659546
Epochs: 8/10 ||| with Training Loss = 0.25803104043006897 ||| Validation Loss =
0.25548702478408813
Epochs: 9/10 ||| with Training Loss = 0.2523624002933502 ||| Validation Loss =
0.25267475843429565
Epochs: 10/10 ||| with Training Loss = 0.24958175420761108 ||| Validation Loss =
0.24853822588920593
```



```
[84]: visualise_model(model7, data_ind, data, model_flag = 1)
```

```
for digit 0  
for digit 1  
for digit 2  
for digit 3  
for digit 4  
for digit 5  
for digit 6  
for digit 7  
for digit 8  
for digit 9
```



7.1 Observations:

- We observe that for unpooling + deconvolution we get the best output for convolutional autoencoders , succeeded by unpooling and deconvolution.
- Convolutional autoencoders take more time to run though, as compared to the standard autoencoders